

Problem 1):

Solution a):

$$c = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Solution b):

$$c = \begin{bmatrix} -1 & 0 & -1 \\ -1 & 0 & -1 \\ -1 & 0 & -1 \end{bmatrix}$$

Solution c):

$$c = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

Solution d):

$$c = \begin{bmatrix} -1 & -1 & 0 \\ -1 & 0 & -1 \\ 0 & -1 & -1 \end{bmatrix}$$

Solution e):

One example of an image operation that cannot be implemented using a 3x3 convolutional filter is a non-linear edge detector. A non-linear edge detector aims to highlight the edges in an image using non-linear operations. This can be achieved using techniques such as morphological operations or non-linear filtering, which involve operations such as dilation, erosion, and median filtering. These non-linear operations cannot be implemented using a 3x3 convolutional filter because a convolutional filter is a linear operator that operates on a local neighborhood of pixels in a fixed way. It applies the same linear transformation to each pixel in the neighborhood, regardless of its value or position. Therefore, a 3x3 convolutional filter cannot capture the complex non-linear relationships between pixels that are necessary for non-linear edge detection.

Problem 2):**Solution a):**

The formula for l_2 loss is given by

$$L(w) = \frac{1}{2} \|y - wx\|^2$$

So the new l_2 loss with λ parameter for regularization is

$$\bar{L}(w) = L(w) + \lambda \|w\|_2^2$$

Solution b):

If the learning rate η , then the general gradient update rule is,

$$w_{t+1} = w_t - \eta \nabla \bar{L}(w)$$

$$w_{t+1} = w_t - \lambda \nabla L(w_t) - 2\eta \lambda w_t$$

$$w_{t+1} = (1 - 2\eta\lambda) w_t - \eta \nabla L(w_t)$$

Solution c):

From the equation mentioned in Solution B, we can see that the updated w consists of shrinking/decaying gradient by a factor of $(1 - 2\eta\lambda)$ and then updating in the direction of the gradient

Solution d):

Increasing λ penalizes the l_2 norm of the weight vector, thus enforcing smaller weights on average. In order for the gradient to be stable, the constraining factor should be smaller than 1, i.e. $\eta < \frac{1}{2\lambda}$

Problem 3):**Solution a):**

The definition of IOU for any two bounding boxes A and B is given by:

$$IOU(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

Since the RHS is non-negative, the number has to be bigger than or equal to 0. Moreover, $A \cap B \subseteq A \cup B$ and hence the numerator has to be no bigger than the denominator. Therefore IOU is bounded between 0 and 1 (inclusive).


Solution b):

Consider two identical size square boxes A and B, both aligned at the same horizontal level. Fix B and then imagine sliding A from left to right. As A moves, the IOU will start from 0, increase until perfect overlap and then decrease until no overlap. The graph we will get is a step function i.e., it jumps from 0 to 1 when the boxes overlap and stays at 0 otherwise, the change in IOU will be discontinuous and will not have a well-defined derivative.

AlexNet

In this problem, you are asked to train a deep convolutional neural network to perform image classification. In fact, this is a slight variation of a network called *AlexNet*. This is a landmark model in deep learning, and arguably kickstarted the current (and ongoing, and massive) wave of innovation in modern AI when its results were first presented in 2012. AlexNet was the first real-world demonstration of a *deep* classifier that was trained end-to-end on data and that outperformed all other ML models thus far.

We will train AlexNet using the [CIFAR10](#) dataset, which consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. The classes are: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck.



A lot of the code you will need is already provided in this notebook; all you need to do is to fill in the missing pieces, and interpret your results.

Warning : AlexNet takes a good amount of time to train (~1 minute per epoch on Google Colab). So please budget enough time to do this homework.

```
In [ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.optim.lr_scheduler import _LRScheduler
import torch.utils.data as data

import torchvision.transforms as transforms
import torchvision.datasets as datasets

from sklearn import decomposition
from sklearn import manifold
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
import matplotlib.pyplot as plt
import numpy as np

import copy
import random
import time
```

```
In [ ]: SEED = 1234

random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.backends.cudnn.deterministic = True
```

Loading and Preparing the Data

Our dataset is made up of color images but three color channels (red, green and blue), compared to MNIST's black and white images with a single color channel. To normalize our data we need to calculate the means and standard deviations for each of the color channels independently, and normalize them.

```
In [ ]: ROOT = '.data'
train_data = datasets.CIFAR10(root = ROOT,
                              train = True,
                              download = True)
```

Files already downloaded and verified

```
In [ ]: # Compute means and standard deviations along the R,G,B channel

means = train_data.data.mean(axis = (0,1,2)) / 255
stds = train_data.data.std(axis = (0,1,2)) / 255
```

Next, we will do data augmentation. For each training image we will randomly rotate it (by up to 5 degrees), flip/mirror with probability 0.5, shift by +/-1 pixel. Finally we will normalize each color channel using the means/stds we calculated above.

```
In [ ]: train_transforms = transforms.Compose([
    transforms.RandomRotation(5),
    transforms.RandomHorizontalFlip(0.5),
    transforms.RandomCrop(32, padding = 2),
    transforms.ToTensor(),
    transforms.Normalize(mean = means,
                        std = stds)
])

test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean = means,
                        std = stds)
])
```

Next, we'll load the dataset along with the transforms defined above.

We will also create a validation set with 10% of the training samples. The validation set will be used to monitor loss along different epochs, and we will pick the model along the optimization path that performed the best, and report final test accuracy numbers using this model.

```
In [ ]: train_data = datasets.CIFAR10(ROOT,
                                     train = True,
                                     download = True,
                                     transform = train_transforms)

test_data = datasets.CIFAR10(ROOT,
                              train = False,
                              download = True,
                              transform = test_transforms)
```

Files already downloaded and verified
Files already downloaded and verified

```
In [ ]: VALID_RATIO = 0.9

n_train_examples = int(len(train_data) * VALID_RATIO)
n_valid_examples = len(train_data) - n_train_examples

train_data, valid_data = data.random_split(train_data,
                                           [n_train_examples, n_valid_examples])
```

```
In [ ]: valid_data = copy.deepcopy(valid_data)
valid_data.dataset.transform = test_transforms
```

Now, we'll create a function to plot some of the images in our dataset to see what they actually look like.

Note that by default PyTorch handles images that are arranged `[channel, height, width]`, but `matplotlib` expects images to be `[height, width, channel]`, hence we need to permute the dimensions of our images before plotting them.

```
In [ ]: def plot_images(images, labels, classes, normalize = False):

    n_images = len(images)

    rows = int(np.sqrt(n_images))
    cols = int(np.sqrt(n_images))

    fig = plt.figure(figsize = (10, 10))

    for i in range(rows*cols):

        ax = fig.add_subplot(rows, cols, i+1)

        image = images[i]

        if normalize:
            image_min = image.min()
            image_max = image.max()
            image.clamp_(min = image_min, max = image_max)
            image.add_(-image_min).div_(image_max - image_min + 1e-5)

        ax.imshow(image.permute(1, 2, 0).cpu().numpy())
        ax.set_title(classes[labels[i]])
        ax.axis('off')
```

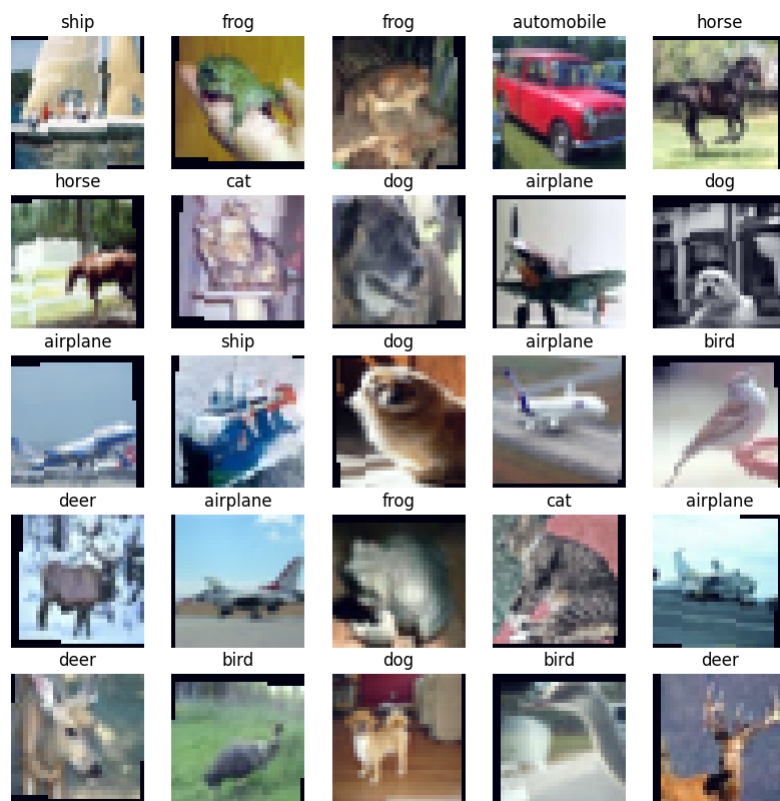
One point here: `matplotlib` is expecting the values of every pixel to be between `[0, 1]`, however our normalization will cause them to be outside this range. By default `matplotlib` will then clip these values into the `[0, 1]` range. This clipping causes all of the images to look a bit weird - all of the colors are oversaturated. The solution is to normalize each image between `[0, 1]`.

```
In [ ]: N_IMAGES = 25

images, labels = zip(*[(image, label) for image, label in
                       [train_data[i] for i in range(N_IMAGES)]])

classes = test_data.classes
```

```
In [ ]: plot_images(images, labels, classes, normalize = True)
```



We'll be normalizing our images by default from now on, so we'll write a function that does it for us which we can use whenever we need to renormalize an image.

```
In [ ]: def normalize_image(image):
        image_min = image.min()
        image_max = image.max()
        image.clamp_(min = image_min, max = image_max)
        image.add_(-image_min).div_(image_max - image_min + 1e-5)
        return image
```

The final bit of the data processing is creating the iterators. We will use a large. Generally, a larger batch size means that our model trains faster but is a bit more susceptible to overfitting.

```
In [ ]: # Q1: Create data loaders for train_data, valid_data, test_data
        # Use batch size 256

        # BATCH_SIZE =
        # train_iterator =
        # valid_iterator =
        # test_iterator =

        BATCH_SIZE = 256
        train_iterator = data.DataLoader(train_data, batch_size=BATCH_SIZE, shuffle=True)
        valid_iterator = data.DataLoader(valid_data, batch_size=BATCH_SIZE)
        test_iterator = data.DataLoader(test_data, batch_size=BATCH_SIZE)
```

Defining the Model

Next up is defining the model.

AlexNet will have the following architecture:

- There are 5 2D convolutional layers (which serve as *feature extractors*), followed by 3 linear layers (which serve as the *classifier*).
- All layers (except the last one) have `ReLU` activations. (Use `inplace=True` while defining your ReLUs.)
- All convolutional filter sizes have kernel size 3×3 and padding 1.
- Convolutional layer 1 has stride 2. All others have the default stride (1).
- Convolutional layers 1, 2, and 5 are followed by a 2D maxpool of size 2.
- Linear layers 1 and 2 are preceded by Dropouts with Bernoulli parameter 0.5.
- For the convolutional layers, the number of channels is set as follows. We start with 3 channels and then proceed like this:

▪ $3 \rightarrow 64 \rightarrow 192 \rightarrow 384 \rightarrow 256 \rightarrow 256$

In the end, if everything is correct you should get a feature map of size $2 \times 2 \times 256 = 1024$.

- For the linear layers, the feature sizes are as follows:

▪ $1024 \rightarrow 4096 \rightarrow 4096 \rightarrow 10$.

(The 10, of course, is because 10 is the number of classes in CIFAR-10).

```
In [ ]: # class AlexNet(nn.Module):
        # def __init__(self, output_dim):
        #     super().__init__()
        #
        #     self.features = nn.Sequential(
```

```

#         # Define according to the steps described above
#         nn.Conv2d(3, 64, kernel_size=11, stride=4, padding=2),
#         nn.ReLU(inplace=True),
#         nn.MaxPool2d(kernel_size=3, stride=2),
#         nn.Conv2d(64, 192, kernel_size=5, padding=2),
#         nn.ReLU(inplace=True),
#         nn.MaxPool2d(kernel_size=3, stride=2),
#         nn.Conv2d(192, 384, kernel_size=3, padding=1),
#         nn.ReLU(inplace=True),
#         nn.Conv2d(384, 256, kernel_size=3, padding=1),
#         nn.ReLU(inplace=True),
#         nn.Conv2d(256, 256, kernel_size=3, padding=1),
#         nn.ReLU(inplace=True),
#         nn.MaxPool2d(kernel_size=3, stride=2)
#     )
#     self.avgpool = nn.AdaptiveAvgPool2d((6, 6))

#     self.classifier = nn.Sequential(
#         # define according to the steps described above
#         nn.Dropout(),
#         nn.Linear(256 * 6 * 6, 4096),
#         nn.ReLU(inplace=True),
#         nn.Dropout(),
#         nn.Linear(4096, 4096),
#         nn.ReLU(inplace=True),
#         nn.Linear(4096, output_dim)
#     )

#     def forward(self, x):
#         x = self.features(x)
#         x = self.avgpool(x)
#         # h = x.view(x.shape[0], -1)
#         h = x.view(x.size(0), 256 * 6 * 6)
#         x = self.classifier(h)
#         return x, h

```

```

In [ ]: class AlexNet(nn.Module):
        def __init__(self, num_classes=10):
            super(AlexNet, self).__init__()
            self.features = nn.Sequential(
                nn.Conv2d(3, 64, kernel_size=3, stride=2, padding=1),
                nn.ReLU(inplace=True),
                nn.MaxPool2d(kernel_size=2),
                nn.Conv2d(64, 192, kernel_size=3, padding=1),
                nn.ReLU(inplace=True),
                nn.MaxPool2d(kernel_size=2),
                nn.Conv2d(192, 384, kernel_size=3, padding=1),
                nn.ReLU(inplace=True),
                nn.Conv2d(384, 256, kernel_size=3, padding=1),
                nn.ReLU(inplace=True),
                nn.Conv2d(256, 256, kernel_size=3, padding=1),
                nn.ReLU(inplace=True),
                nn.MaxPool2d(kernel_size=2),
            )
            self.avgpool = nn.AdaptiveAvgPool2d((2, 2))
            self.classifier = nn.Sequential(
                nn.Dropout(),
                nn.Linear(256 * 2 * 2, 4096),
                nn.ReLU(inplace=True),
                nn.Dropout(),
                nn.Linear(4096, 4096),
                nn.ReLU(inplace=True),
                nn.Linear(4096, num_classes),
            )

        def forward(self, x):
            x = self.features(x)
            x = self.avgpool(x)
            h = x.view(x.size(0), 256 * 2 * 2)
            x = self.classifier(h)
            return x, h

```

We'll create an instance of our model with the desired amount of classes.

```

In [ ]: OUTPUT_DIM = 10
        model = AlexNet(OUTPUT_DIM)

```

Training the Model

We first initialize parameters in PyTorch by creating a function that takes in a PyTorch module, checking what type of module it is, and then using the `nn.init` methods to actually initialize the parameters.

For convolutional layers we will initialize using the *Kaiming Normal* scheme, also known as *He Normal*. For the linear layers we initialize using the *Xavier Normal* scheme, also known as *Glorot Normal*. For both types of layer we initialize the bias terms to zeros.

```

In [ ]: def initialize_parameters(m):
        if isinstance(m, nn.Conv2d):
            nn.init.kaiming_normal_(m.weight.data, nonlinearity = 'relu')
            nn.init.constant_(m.bias.data, 0)
        elif isinstance(m, nn.Linear):
            nn.init.xavier_normal_(m.weight.data, gain = nn.init.calculate_gain('relu'))
            nn.init.constant_(m.bias.data, 0)

```

We apply the initialization by using the model's `apply` method. If your definitions above are correct you should get the printed output as below.

```

In [ ]: model.apply(initialize_parameters)

```

```

Out[ ]: AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (6): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(2, 2))
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=1024, out_features=4096, bias=True)
    (2): ReLU(inplace=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=4096, out_features=4096, bias=True)
    (5): ReLU(inplace=True)
    (6): Linear(in_features=4096, out_features=10, bias=True)
  )
)

```

We then define the loss function we want to use, the device we'll use and place our model and criterion on to our device.

```

In [ ]: optimizer = optim.Adam(model.parameters(), lr = 1e-3)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
criterion = nn.CrossEntropyLoss()

model = model.to(device)
criterion = criterion.to(device)

```

This is formatted as code

We define a function to calculate accuracy...

```

In [ ]: def calculate_accuracy(y_pred, y):
    top_pred = y_pred.argmax(1, keepdim = True)
    correct = top_pred.eq(y.view_as(top_pred)).sum()
    acc = correct.float() / y.shape[0]
    return acc

```

As we are using dropout we need to make sure to "turn it on" when training by using `model.train()`.

```

In [ ]: def train(model, iterator, optimizer, criterion, device):

    epoch_loss = 0
    epoch_acc = 0

    model.train()

    for (x, y) in iterator:

        x = x.to(device)
        y = y.to(device)

        optimizer.zero_grad()

        y_pred, _ = model(x)

        loss = criterion(y_pred, y)

        acc = calculate_accuracy(y_pred, y)

        loss.backward()

        optimizer.step()

        epoch_loss += loss.item()
        epoch_acc += acc.item()

    return epoch_loss / len(iterator), epoch_acc / len(iterator)

```

We also define an evaluation loop, making sure to "turn off" dropout with `model.eval()`.

```

In [ ]: def evaluate(model, iterator, criterion, device):

    epoch_loss = 0
    epoch_acc = 0

    model.eval()

    with torch.no_grad():

        for (x, y) in iterator:

            x = x.to(device)
            y = y.to(device)

            y_pred, _ = model(x)

            loss = criterion(y_pred, y)

            acc = calculate_accuracy(y_pred, y)

            epoch_loss += loss.item()
            epoch_acc += acc.item()

    return epoch_loss / len(iterator), epoch_acc / len(iterator)

```

Next, we define a function to tell us how long an epoch takes.

```

In [ ]: def epoch_time(start_time, end_time):
    elapsed_time = end_time - start_time
    elapsed_mins = int(elapsed_time / 60)

```



```

    elapsed_secs = int(elapsed_time - (elapsed_mins * 60))
    return elapsed_mins, elapsed_secs

```

Then, finally, we train our model.

Train it for 25 epochs (using the train dataset). At the end of each epoch, compute the validation loss and keep track of the best model. You might find the command `torch.save` helpful.

At the end you should expect to see validation losses of ~76% accuracy.

```

In [ ]: # Q3: train your model here for 25 epochs.
        # Print out training and validation loss/accuracy of the model after each epoch
        # Keep track of the model that achieved best validation loss thus far.

EPOCHS = 25

# Fill training code here

best_valid_loss = float('inf')
for epoch in range(EPOCHS):
    train_loss, train_acc = train(model, train_iterator, optimizer, criterion, device)
    valid_loss, valid_acc = evaluate(model, valid_iterator, criterion, device)

    if valid_loss < best_valid_loss:
        best_valid_loss = valid_loss
        torch.save(model.state_dict(), 'best_model.pt')

    print(f'Epoch: {epoch+1:02}')
    print(f'\tTrain Loss: {train_loss:.3f} | Train Acc: {train_acc*100:.2f}%')
    print(f'\tVal. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*100:.2f}%')

```

```

Epoch: 01
  Train Loss: 2.385 | Train Acc: 21.71%
  Val. Loss: 1.613 | Val. Acc: 39.04%
Epoch: 02
  Train Loss: 1.539 | Train Acc: 43.00%
  Val. Loss: 1.351 | Val. Acc: 50.60%
Epoch: 03
  Train Loss: 1.357 | Train Acc: 50.75%
  Val. Loss: 1.241 | Val. Acc: 54.72%
Epoch: 04
  Train Loss: 1.256 | Train Acc: 54.90%
  Val. Loss: 1.141 | Val. Acc: 58.73%
Epoch: 05
  Train Loss: 1.173 | Train Acc: 58.41%
  Val. Loss: 1.083 | Val. Acc: 61.63%
Epoch: 06
  Train Loss: 1.112 | Train Acc: 60.57%
  Val. Loss: 1.059 | Val. Acc: 62.47%
Epoch: 07
  Train Loss: 1.066 | Train Acc: 62.07%
  Val. Loss: 1.018 | Val. Acc: 65.02%
Epoch: 08
  Train Loss: 1.008 | Train Acc: 64.62%
  Val. Loss: 0.997 | Val. Acc: 64.91%
Epoch: 09
  Train Loss: 0.966 | Train Acc: 66.29%
  Val. Loss: 0.936 | Val. Acc: 67.48%
Epoch: 10
  Train Loss: 0.929 | Train Acc: 67.48%
  Val. Loss: 0.919 | Val. Acc: 68.54%
Epoch: 11
  Train Loss: 0.899 | Train Acc: 68.67%
  Val. Loss: 0.922 | Val. Acc: 67.75%
Epoch: 12
  Train Loss: 0.876 | Train Acc: 69.48%
  Val. Loss: 0.835 | Val. Acc: 71.28%
Epoch: 13
  Train Loss: 0.841 | Train Acc: 70.71%
  Val. Loss: 0.849 | Val. Acc: 70.74%
Epoch: 14
  Train Loss: 0.827 | Train Acc: 71.38%
  Val. Loss: 0.832 | Val. Acc: 72.24%
Epoch: 15
  Train Loss: 0.800 | Train Acc: 72.07%
  Val. Loss: 0.804 | Val. Acc: 72.87%
Epoch: 16
  Train Loss: 0.770 | Train Acc: 73.30%
  Val. Loss: 0.802 | Val. Acc: 72.40%
Epoch: 17
  Train Loss: 0.756 | Train Acc: 73.56%
  Val. Loss: 0.784 | Val. Acc: 73.14%
Epoch: 18
  Train Loss: 0.739 | Train Acc: 74.43%
  Val. Loss: 0.753 | Val. Acc: 74.50%
Epoch: 19
  Train Loss: 0.721 | Train Acc: 75.17%
  Val. Loss: 0.767 | Val. Acc: 73.64%
Epoch: 20
  Train Loss: 0.710 | Train Acc: 75.37%
  Val. Loss: 0.754 | Val. Acc: 74.73%
Epoch: 21
  Train Loss: 0.692 | Train Acc: 76.04%
  Val. Loss: 0.757 | Val. Acc: 74.90%
Epoch: 22
  Train Loss: 0.679 | Train Acc: 76.66%
  Val. Loss: 0.748 | Val. Acc: 75.09%
Epoch: 23
  Train Loss: 0.673 | Train Acc: 76.63%
  Val. Loss: 0.727 | Val. Acc: 75.59%
Epoch: 24
  Train Loss: 0.650 | Train Acc: 77.80%
  Val. Loss: 0.726 | Val. Acc: 75.35%
Epoch: 25
  Train Loss: 0.644 | Train Acc: 78.15%
  Val. Loss: 0.756 | Val. Acc: 74.38%

```

Evaluating the model

We then load the parameters of our model that achieved the best validation loss. You should expect to see ~75% accuracy of this model on the test dataset.

Finally, plot the confusion matrix of this model and comment on any interesting patterns you can observe there. For example, which two classes are confused the most?

```
In [ ]: # Q4: Load the best performing model, evaluate it on the test dataset, and print test accuracy.
# Also, print out the confusion matrix.
```


```
In [ ]: # def get_predictions(model, iterator, device):
#     model.eval()
#     labels = []
#     probs = []
#     # Q4: Fill code here.
#     labels = torch.cat(labels, dim = 0)
#     probs = torch.cat(probs, dim = 0)
#     return labels, probs
def get_predictions(model, iterator, device):
    model.eval()
    labels = []
    probs = []
    with torch.no_grad():
        for batch in iterator:
            x, y = batch
            x = x.to(device)
            y = y.to(device)
            y_pred, _ = model(x)
            labels.append(y)
            probs.append(torch.softmax(y_pred, dim=1))
    labels = torch.cat(labels, dim=0)
    probs = torch.cat(probs, dim=0)
    return labels, probs
```

```
In [ ]: labels, probs = get_predictions(model, test_iterator, device)
```

```
In [ ]: pred_labels = torch.argmax(probs, 1)
```

```
In [ ]: def plot_confusion_matrix(labels, pred_labels, classes):
    fig = plt.figure(figsize = (10, 10));
    ax = fig.add_subplot(1, 1, 1);
    cm = confusion_matrix(labels, pred_labels);
    cm = ConfusionMatrixDisplay(cm, display_labels = classes);
    cm.plot(values_format = 'd', cmap = 'Blues', ax = ax)
    plt.xticks(rotation = 20)
```

```
In [ ]: plot_confusion_matrix(labels, pred_labels, classes)
```



Conclusion

That's it! As a side project (this is not for credit and won't be graded), feel free to play around with different design choices that you made while building this network.

- Whether or not to normalize the color channels in the input.
- The learning rate parameter in Adam.
- The batch size.
- The number of training epochs.

- (and if you are feeling brave -- the AlexNet architecture itself.)

```

In [8]: import os
import numpy as np
import torch
from PIL import Image

class PennFudanDataset(torch.utils.data.Dataset):
    def __init__(self, root, transforms):
        self.root = root
        self.transforms = transforms
        # load all image files, sorting them to
        # ensure that they are aligned
        self.imgs = list(sorted(os.listdir(os.path.join(root, "PNGImages"))))
        self.masks = list(sorted(os.listdir(os.path.join(root, "PedMasks"))))

    def __getitem__(self, idx):
        # load images and masks
        img_path = os.path.join(self.root, "PNGImages", self.imgs[idx])
        mask_path = os.path.join(self.root, "PedMasks", self.masks[idx])
        img = Image.open(img_path).convert("RGB")
        # note that we haven't converted the mask to RGB,
        # because each color corresponds to a different instance
        # with 0 being background
        mask = Image.open(mask_path)
        # convert the PIL Image into a numpy array
        mask = np.array(mask)
        # instances are encoded as different colors
        obj_ids = np.unique(mask)
        # first id is the background, so remove it
        obj_ids = obj_ids[1:]

        # split the color-encoded mask into a set
        # of binary masks
        masks = mask == obj_ids[:, None, None]

        # get bounding box coordinates for each mask
        num_objs = len(obj_ids)
        boxes = []
        for i in range(num_objs):
            pos = np.where(masks[i])
            xmin = np.min(pos[1])
            xmax = np.max(pos[1])
            ymin = np.min(pos[0])
            ymax = np.max(pos[0])
            boxes.append([xmin, ymin, xmax, ymax])

        # convert everything into a torch.Tensor
        boxes = torch.as_tensor(boxes, dtype=torch.float32)
        # there is only one class
        labels = torch.ones((num_objs, ), dtype=torch.int64)
        masks = torch.as_tensor(masks, dtype=torch.uint8)

        image_id = torch.tensor([idx])
        area = (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:, 0])
        # suppose all instances are not crowd
        iscrowd = torch.zeros((num_objs, ), dtype=torch.int64)

        target = {}
        target["boxes"] = boxes
        target["labels"] = labels
        target["masks"] = masks
        target["image_id"] = image_id
        target["area"] = area
        target["iscrowd"] = iscrowd

        if self.transforms is not None:
            img, target = self.transforms(img, target)

        return img, target

    def __len__(self):
        return len(self.imgs)

```

```

In [9]: import torchvision
import torchvision
from torchvision.models.detection import FasterRCNN
from torchvision.models.detection.rpn import AnchorGenerator
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor

def get_model_instance_segmentation_option1(num_classes):
    # load an instance segmentation model pre-trained pre-trained on COCO
    model = torchvision.models.detection.maskrcnn_resnet50_fpn(pretrained=True)

    # get number of input features for the classifier
    in_features = model.roi_heads.box_predictor.cls_score.in_features

    # replace the pre-trained head with a new one
    model.roi_heads.box_predictor = FastRCNNPredictor(in_features, num_classes)

    # now get the number of input features for the mask classifier
    in_features_mask = model.roi_heads.mask_predictor.conv5_mask.in_channels
    hidden_layer = 256

    # and replace the mask predictor with a new one
    model.roi_heads.mask_predictor = MaskRCNNPredictor(in_features_mask,
                                                         hidden_layer,
                                                         num_classes)

    return model

def get_model_instance_segmentation_option2(num_classes):
    # load a pre-trained model for classification and return
    # only the features
    backbone = torchvision.models.mobilenet_v2(weights="DEFAULT").features

    # FasterRCNN needs to know the number of
    # output channels in a backbone. For mobilenet_v2, it's 1280
    # so we need to add it here
    backbone.out_channels = 1280

```

```

# let's make the RPN generate 5 x 3 anchors per spatial
# location, with 5 different sizes and 3 different aspect
# ratios. We have a Tuple[Tuple[int]] because each feature
# map could potentially have different sizes and
# aspect ratios
anchor_generator = AnchorGenerator(sizes=((32, 64, 128, 256, 512).),
                                  aspect_ratios=((0.5, 1.0, 2.0),))

# let's define what are the feature maps that we will
# use to perform the region of interest cropping, as well as
# the size of the crop after rescaling.
# if your backbone returns a Tensor, featmap_names is expected to
# be [0]. More generally, the backbone should return an
# OrderedDict[Tensor], and in featmap_names you can choose which
# feature maps to use.
roi_pooler = torchvision.ops.MultiScaleRoIAlign(featmap_names=['0'],
                                                output_size=7,
                                                sampling_ratio=2)

# put the pieces together inside a FasterRCNN model
model = FasterRCNN(backbone,
                   num_classes=2,
                   rpn_anchor_generator=anchor_generator,
                   box_roi_pool=roi_pooler)

return model

```

In [10]: `import transforms as T`

```

def get_transform(train):
    transforms = []
    transforms.append(T.PILToTensor())
    transforms.append(T.ConvertImageDtype(torch.float))
    if train:
        transforms.append(T.RandomHorizontalFlip(0.5))
    return T.Compose(transforms)

```

In [11]: `from engine import train_one_epoch, evaluate`

```

import utils
from PIL import Image
import torch
import torchvision.transforms as transforms
import requests
import utils
#%matplotlib inline
import matplotlib.pyplot as plt

# train on the GPU or on the CPU, if a GPU is not available
device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')

# our dataset has two classes only - background and person
num_classes = 2

# use our dataset and defined transformations
dataset = PennFudanDataset('/content/drive/MyDrive/Colab Notebooks/PennFudanPed', get_transform(train=True))
dataset_test = PennFudanDataset('/content/drive/MyDrive/Colab Notebooks/PennFudanPed', get_transform(train=False))

# split the dataset in train and test set
indices = torch.randperm(len(dataset)).tolist()
dataset = torch.utils.data.Subset(dataset, indices[:-50])
dataset_test = torch.utils.data.Subset(dataset_test, indices[-50:])

# define training and validation data loaders
data_loader = torch.utils.data.DataLoader(
    dataset, batch_size=2, shuffle=True, num_workers=0,
    collate_fn=utils.collate_fn)

data_loader_test = torch.utils.data.DataLoader(
    dataset_test, batch_size=1, shuffle=False, num_workers=0,
    collate_fn=utils.collate_fn)

```

In [12]: `from google.colab import drive`
`drive.mount('/content/drive')`

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

In [13]: `#MODEL_OPTION1`

```

# get the model using our helper function
model_option1 = get_model_instance_segmentation_option1(num_classes)

# move model to the right device
model_option1.to(device)

# construct an optimizer for option1
params = [p for p in model_option1.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(params, lr=0.005,
                             momentum=0.9, weight_decay=0.0005)

# and a learning rate scheduler for the same
lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
                                                step_size=3,
                                                gamma=0.1)

# let's train it for 10 epochs
num_epochs = 10
print ("MODEL OPTION 1")

for epoch in range(num_epochs):
    print("Model 1 Epoch: {}".format(epoch))
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model_option1, optimizer, data_loader, device, epoch, print_freq=10)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the test dataset
    evaluate(model_option1, data_loader_test, device=device)

print ("\n")
print("That's it! Model 1")

```

```
/usr/local/lib/python3.9/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.
  warnings.warn(
/usr/local/lib/python3.9/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or 'None' for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=MaskRCNN_ResNet50_FPN_Weights.COCO_V1`. You can also use `weights=MaskRCNN_ResNet50_FPN_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)
```

```

MODEL OPTION 1
Model 1 Epoch: 0
Epoch: [0] [ 0/60] eta: 0:01:56 lr: 0.000090 loss: 5.8660 (5.8660) loss_classifier: 0.7740 (0.7740) loss_box_reg: 0.3408 (0.3408) loss_mask: 4.7247 (4.7247) loss_ob
jctness: 0.0241 (0.0241) loss_rpn_box_reg: 0.0023 (0.0023) time: 1.9364 data: 1.3563 max mem: 2579
Epoch: [0] [10/60] eta: 0:01:33 lr: 0.000936 loss: 2.3789 (3.1197) loss_classifier: 0.5083 (0.4910) loss_box_reg: 0.3408 (0.2984) loss_mask: 1.4818 (2.3071) loss_ob
jctness: 0.0192 (0.0185) loss_rpn_box_reg: 0.0028 (0.0047) time: 1.8733 data: 1.3209 max mem: 3528
Epoch: [0] [20/60] eta: 0:01:14 lr: 0.001783 loss: 0.9242 (2.0173) loss_classifier: 0.2323 (0.3516) loss_box_reg: 0.3054 (0.2896) loss_mask: 0.4465 (1.3515) loss_ob
jctness: 0.0192 (0.0187) loss_rpn_box_reg: 0.0030 (0.0059) time: 1.8659 data: 1.3193 max mem: 3528
Epoch: [0] [30/60] eta: 0:00:56 lr: 0.002629 loss: 0.5756 (1.5340) loss_classifier: 0.0937 (0.2659) loss_box_reg: 0.2108 (0.2590) loss_mask: 0.2391 (0.9873) loss_ob
jctness: 0.0099 (0.0160) loss_rpn_box_reg: 0.0059 (0.0058) time: 1.9024 data: 1.3500 max mem: 3554
Epoch: [0] [40/60] eta: 0:00:39 lr: 0.003476 loss: 0.4931 (1.2783) loss_classifier: 0.0713 (0.2166) loss_box_reg: 0.2089 (0.2468) loss_mask: 0.2166 (0.7948) loss_ob
jctness: 0.0081 (0.0141) loss_rpn_box_reg: 0.0059 (0.0059) time: 2.0443 data: 1.4791 max mem: 3554
Epoch: [0] [50/60] eta: 0:00:19 lr: 0.004323 loss: 0.4307 (1.1157) loss_classifier: 0.0552 (0.1853) loss_box_reg: 0.1454 (0.2312) loss_mask: 0.1952 (0.6806) loss_ob
jctness: 0.0045 (0.0122) loss_rpn_box_reg: 0.0057 (0.0065) time: 2.0565 data: 1.4892 max mem: 3559
Epoch: [0] [59/60] eta: 0:00:01 lr: 0.005000 loss: 0.3701 (0.9989) loss_classifier: 0.0367 (0.1630) loss_box_reg: 0.1329 (0.2163) loss_mask: 0.1797 (0.6026) loss_ob
jctness: 0.0016 (0.0106) loss_rpn_box_reg: 0.0056 (0.0063) time: 1.9448 data: 1.3731 max mem: 3742
Epoch: [0] Total time: 0:01:57 (1.9522 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:09 model_time: 0.1674 (0.1674) evaluator_time: 0.0149 (0.0149) time: 0.1996 data: 0.0165 max mem: 3742
Test: [49/50] eta: 0:00:00 model_time: 0.1023 (0.1111) evaluator_time: 0.0067 (0.0120) time: 0.1457 data: 0.0200 max mem: 3742
Test: Total time: 0:00:07 (0.1417 s / it)
Averaged stats: model_time: 0.1023 (0.1111) evaluator_time: 0.0067 (0.0120)
Accumulating evaluation results...
DONE (t=0.03s).
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.593
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.968
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.766
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.381
Average Precision (AP) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.499
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.604
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.260
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.660
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.660
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.500
Average Recall (AR) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.633
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.666
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.611
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.972
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.762
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.330
Average Precision (AP) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.337
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.626
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.272
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.658
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.663
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.533
Average Recall (AR) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.678
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.665
Model 1 Epoch: 1
Epoch: [1] [ 0/60] eta: 0:00:35 lr: 0.005000 loss: 0.4237 (0.4237) loss_classifier: 0.0566 (0.0566) loss_box_reg: 0.1427 (0.1427) loss_mask: 0.2189 (0.2189) loss_ob
jctness: 0.0003 (0.0003) loss_rpn_box_reg: 0.0052 (0.0052) time: 0.5921 data: 0.0361 max mem: 3742
Epoch: [1] [10/60] eta: 0:00:31 lr: 0.005000 loss: 0.3257 (0.3524) loss_classifier: 0.0435 (0.0462) loss_box_reg: 0.1320 (0.1303) loss_mask: 0.1620 (0.1690) loss_ob
jctness: 0.0017 (0.0024) loss_rpn_box_reg: 0.0031 (0.0045) time: 0.6349 data: 0.0409 max mem: 3742
Epoch: [1] [20/60] eta: 0:00:25 lr: 0.005000 loss: 0.3009 (0.3246) loss_classifier: 0.0401 (0.0439) loss_box_reg: 0.1076 (0.1142) loss_mask: 0.1506 (0.1588) loss_ob
jctness: 0.0018 (0.0025) loss_rpn_box_reg: 0.0040 (0.0053) time: 0.6325 data: 0.0400 max mem: 3742
Epoch: [1] [30/60] eta: 0:00:18 lr: 0.005000 loss: 0.2838 (0.3094) loss_classifier: 0.0400 (0.0424) loss_box_reg: 0.0880 (0.1051) loss_mask: 0.1452 (0.1543) loss_ob
jctness: 0.0012 (0.0023) loss_rpn_box_reg: 0.0053 (0.0054) time: 0.6196 data: 0.0378 max mem: 3742
Epoch: [1] [40/60] eta: 0:00:12 lr: 0.005000 loss: 0.2544 (0.2860) loss_classifier: 0.0311 (0.0379) loss_box_reg: 0.0549 (0.0915) loss_mask: 0.1316 (0.1493) loss_ob
jctness: 0.0008 (0.0021) loss_rpn_box_reg: 0.0033 (0.0052) time: 0.6021 data: 0.0364 max mem: 3742
Epoch: [1] [50/60] eta: 0:00:06 lr: 0.005000 loss: 0.2330 (0.2792) loss_classifier: 0.0262 (0.0377) loss_box_reg: 0.0546 (0.0882) loss_mask: 0.1263 (0.1461) loss_ob
jctness: 0.0007 (0.0022) loss_rpn_box_reg: 0.0032 (0.0049) time: 0.6087 data: 0.0352 max mem: 3742
Epoch: [1] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2324 (0.2717) loss_classifier: 0.0275 (0.0368) loss_box_reg: 0.0553 (0.0841) loss_mask: 0.1263 (0.1440) loss_ob
jctness: 0.0004 (0.0020) loss_rpn_box_reg: 0.0032 (0.0049) time: 0.6104 data: 0.0345 max mem: 3742
Epoch: [1] Total time: 0:00:36 (0.6159 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:10 model_time: 0.1657 (0.1657) evaluator_time: 0.0120 (0.0120) time: 0.2010 data: 0.0225 max mem: 3742
Test: [49/50] eta: 0:00:00 model_time: 0.0985 (0.1064) evaluator_time: 0.0037 (0.0055) time: 0.1263 data: 0.0167 max mem: 3742
Test: Total time: 0:00:06 (0.1297 s / it)
Averaged stats: model_time: 0.0985 (0.1064) evaluator_time: 0.0037 (0.0055)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.768
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.982
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.942
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.363
Average Precision (AP) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.618
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.785
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.336
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.808
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.808
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.500
Average Recall (AR) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.733
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.822
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.741
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.982
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.926
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.370
Average Precision (AP) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.506
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.759
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.321
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.770
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.770
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.600
Average Recall (AR) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.689
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.781
Model 1 Epoch: 2
Epoch: [2] [ 0/60] eta: 0:00:43 lr: 0.005000 loss: 0.2319 (0.2319) loss_classifier: 0.0396 (0.0396) loss_box_reg: 0.0734 (0.0734) loss_mask: 0.1158 (0.1158) loss_ob
jctness: 0.0003 (0.0003) loss_rpn_box_reg: 0.0029 (0.0029) time: 0.7233 data: 0.0464 max mem: 3742
Epoch: [2] [10/60] eta: 0:00:30 lr: 0.005000 loss: 0.2233 (0.2445) loss_classifier: 0.0276 (0.0330) loss_box_reg: 0.0562 (0.0700) loss_mask: 0.1351 (0.1349) loss_ob
jctness: 0.0006 (0.0012) loss_rpn_box_reg: 0.0046 (0.0054) time: 0.6173 data: 0.0452 max mem: 3742
Epoch: [2] [20/60] eta: 0:00:23 lr: 0.005000 loss: 0.1867 (0.2278) loss_classifier: 0.0240 (0.0292) loss_box_reg: 0.0483 (0.0615) loss_mask: 0.1240 (0.1324) loss_ob
jctness: 0.0004 (0.0009) loss_rpn_box_reg: 0.0029 (0.0038) time: 0.5888 data: 0.0386 max mem: 3742
Epoch: [2] [30/60] eta: 0:00:18 lr: 0.005000 loss: 0.2387 (0.2389) loss_classifier: 0.0289 (0.0308) loss_box_reg: 0.0599 (0.0674) loss_mask: 0.1259 (0.1353) loss_ob
jctness: 0.0004 (0.0012) loss_rpn_box_reg: 0.0026 (0.0042) time: 0.6107 data: 0.0351 max mem: 3742
Epoch: [2] [40/60] eta: 0:00:12 lr: 0.005000 loss: 0.2450 (0.2316) loss_classifier: 0.0295 (0.0310) loss_box_reg: 0.0661 (0.0635) loss_mask: 0.1324 (0.1319) loss_ob

```

```

jectness: 0.0007 (0.0012) loss_rpn_box_reg: 0.0047 (0.0041) time: 0.6250 data: 0.0382 max mem: 3748
Epoch: [2] [50/60] eta: 0:00:06 lr: 0.005000 loss: 0.2110 (0.2333) loss_classifier: 0.0314 (0.0322) loss_box_reg: 0.0511 (0.0652) loss_mask: 0.1215 (0.1306) loss_ob
jectness: 0.0004 (0.0011) loss_rpn_box_reg: 0.0037 (0.0042) time: 0.5832 data: 0.0360 max mem: 3748
Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2345 (0.2314) loss_classifier: 0.0341 (0.0316) loss_box_reg: 0.0635 (0.0641) loss_mask: 0.1213 (0.1306) loss_ob
jectness: 0.0003 (0.0010) loss_rpn_box_reg: 0.0034 (0.0041) time: 0.5916 data: 0.0374 max mem: 3748
Epoch: [2] Total time: 0:00:36 (0.6053 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:09 model_time: 0.1700 (0.1700) evaluator_time: 0.0067 (0.0067) time: 0.1945 data: 0.0171 max mem: 3748
Test: [49/50] eta: 0:00:00 model_time: 0.0988 (0.1064) evaluator_time: 0.0035 (0.0054) time: 0.1285 data: 0.0173 max mem: 3748
Test: Total time: 0:00:06 (0.1296 s / it)
Averaged stats: model_time: 0.0988 (0.1064) evaluator_time: 0.0035 (0.0054)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.788
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.982
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.945
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.374
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.612
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.805
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.343
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.828
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.828
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.533
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.756
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.842
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.731
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.983
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.897
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.400
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.542
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.744
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.322
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.766
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.768
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.633
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.744
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.773
Model 1 Epoch: 3
Epoch: [3] [ 0/60] eta: 0:00:36 lr: 0.000500 loss: 0.2125 (0.2125) loss_classifier: 0.0267 (0.0267) loss_box_reg: 0.0480 (0.0480) loss_mask: 0.1368 (0.1368) loss_ob
jectness: 0.0001 (0.0001) loss_rpn_box_reg: 0.0009 (0.0009) time: 0.6009 data: 0.0399 max mem: 3748
Epoch: [3] [10/60] eta: 0:00:30 lr: 0.000500 loss: 0.2125 (0.2237) loss_classifier: 0.0285 (0.0297) loss_box_reg: 0.0511 (0.0574) loss_mask: 0.1352 (0.1306) loss_ob
jectness: 0.0006 (0.0014) loss_rpn_box_reg: 0.0041 (0.0047) time: 0.6178 data: 0.0428 max mem: 3748
Epoch: [3] [20/60] eta: 0:00:23 lr: 0.000500 loss: 0.1779 (0.1978) loss_classifier: 0.0258 (0.0269) loss_box_reg: 0.0359 (0.0461) loss_mask: 0.1114 (0.1203) loss_ob
jectness: 0.0005 (0.0010) loss_rpn_box_reg: 0.0027 (0.0035) time: 0.5927 data: 0.0367 max mem: 3748
Epoch: [3] [30/60] eta: 0:00:18 lr: 0.000500 loss: 0.1763 (0.1988) loss_classifier: 0.0246 (0.0268) loss_box_reg: 0.0328 (0.0486) loss_mask: 0.1049 (0.1188) loss_ob
jectness: 0.0005 (0.0010) loss_rpn_box_reg: 0.0024 (0.0036) time: 0.6036 data: 0.0373 max mem: 3748
Epoch: [3] [40/60] eta: 0:00:12 lr: 0.000500 loss: 0.1909 (0.1962) loss_classifier: 0.0260 (0.0266) loss_box_reg: 0.0494 (0.0467) loss_mask: 0.1148 (0.1185) loss_ob
jectness: 0.0004 (0.0010) loss_rpn_box_reg: 0.0030 (0.0033) time: 0.6157 data: 0.0384 max mem: 3748
Epoch: [3] [50/60] eta: 0:00:06 lr: 0.000500 loss: 0.1962 (0.1967) loss_classifier: 0.0260 (0.0270) loss_box_reg: 0.0474 (0.0466) loss_mask: 0.1155 (0.1185) loss_ob
jectness: 0.0004 (0.0011) loss_rpn_box_reg: 0.0021 (0.0035) time: 0.6055 data: 0.0383 max mem: 3748
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1837 (0.1939) loss_classifier: 0.0250 (0.0267) loss_box_reg: 0.0397 (0.0455) loss_mask: 0.1110 (0.1174) loss_ob
jectness: 0.0005 (0.0011) loss_rpn_box_reg: 0.0028 (0.0034) time: 0.6071 data: 0.0379 max mem: 3748
Epoch: [3] Total time: 0:00:36 (0.6063 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:10 model_time: 0.1773 (0.1773) evaluator_time: 0.0087 (0.0087) time: 0.2051 data: 0.0183 max mem: 3748
Test: [49/50] eta: 0:00:00 model_time: 0.0974 (0.1074) evaluator_time: 0.0032 (0.0066) time: 0.1257 data: 0.0172 max mem: 3748
Test: Total time: 0:00:06 (0.1344 s / it)
Averaged stats: model_time: 0.0974 (0.1074) evaluator_time: 0.0032 (0.0066)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.808
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.987
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.949
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.384
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.627
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.826
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.352
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.840
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.840
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.533
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.756
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.855
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.741
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.978
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.878
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.396
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.531
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.755
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.328
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.772
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.772
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.756
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.778
Model 1 Epoch: 4
Epoch: [4] [ 0/60] eta: 0:00:29 lr: 0.000500 loss: 0.1836 (0.1836) loss_classifier: 0.0257 (0.0257) loss_box_reg: 0.0415 (0.0415) loss_mask: 0.1111 (0.1111) loss_ob
jectness: 0.0040 (0.0040) loss_rpn_box_reg: 0.0013 (0.0013) time: 0.4944 data: 0.0278 max mem: 3748
Epoch: [4] [10/60] eta: 0:00:28 lr: 0.000500 loss: 0.1576 (0.1762) loss_classifier: 0.0249 (0.0266) loss_box_reg: 0.0334 (0.0379) loss_mask: 0.1037 (0.1091) loss_ob
jectness: 0.0009 (0.0012) loss_rpn_box_reg: 0.0009 (0.0013) time: 0.5626 data: 0.0343 max mem: 3748
Epoch: [4] [20/60] eta: 0:00:23 lr: 0.000500 loss: 0.1803 (0.1852) loss_classifier: 0.0242 (0.0269) loss_box_reg: 0.0369 (0.0417) loss_mask: 0.1173 (0.1138) loss_ob
jectness: 0.0003 (0.0008) loss_rpn_box_reg: 0.0018 (0.0020) time: 0.5829 data: 0.0385 max mem: 3748
Epoch: [4] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1872 (0.1918) loss_classifier: 0.0242 (0.0262) loss_box_reg: 0.0403 (0.0427) loss_mask: 0.1211 (0.1197) loss_ob
jectness: 0.0003 (0.0008) loss_rpn_box_reg: 0.0024 (0.0025) time: 0.6006 data: 0.0399 max mem: 3748
Epoch: [4] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1794 (0.1886) loss_classifier: 0.0227 (0.0258) loss_box_reg: 0.0361 (0.0418) loss_mask: 0.1128 (0.1175) loss_ob
jectness: 0.0006 (0.0008) loss_rpn_box_reg: 0.0024 (0.0027) time: 0.6085 data: 0.0374 max mem: 3748
Epoch: [4] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1797 (0.1870) loss_classifier: 0.0268 (0.0263) loss_box_reg: 0.0384 (0.0411) loss_mask: 0.1041 (0.1159) loss_ob
jectness: 0.0007 (0.0012) loss_rpn_box_reg: 0.0022 (0.0027) time: 0.6102 data: 0.0356 max mem: 3748
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1908 (0.1876) loss_classifier: 0.0295 (0.0266) loss_box_reg: 0.0431 (0.0418) loss_mask: 0.1082 (0.1150) loss_ob
jectness: 0.0008 (0.0013) loss_rpn_box_reg: 0.0022 (0.0029) time: 0.6307 data: 0.0368 max mem: 3748
Epoch: [4] Total time: 0:00:36 (0.6054 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:09 model_time: 0.1630 (0.1630) evaluator_time: 0.0052 (0.0052) time: 0.1859 data: 0.0170 max mem: 3748
Test: [49/50] eta: 0:00:00 model_time: 0.0977 (0.1057) evaluator_time: 0.0028 (0.0049) time: 0.1254 data: 0.0165 max mem: 3748
Test: Total time: 0:00:06 (0.1282 s / it)

```



```

Averaged stats: model_time: 0.0977 (0.1057) evaluator_time: 0.0028 (0.0049)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.826
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.987
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.950
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.413
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.650
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.843
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.357
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.857
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.857
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.872
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.752
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.987
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.901
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.385
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.523
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.331
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.780
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.780
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.756
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.788
Model 1 Epoch: 5
Epoch: [5] [ 0/60] eta: 0:00:28 lr: 0.000500 loss: 0.1180 (0.1180) loss_classifier: 0.0121 (0.0121) loss_box_reg: 0.0125 (0.0125) loss_mask: 0.0923 (0.0923) loss_ob
jctness: 0.0000 (0.0000) loss_rpn_box_reg: 0.0011 (0.0011) time: 0.4785 data: 0.0294 max mem: 3748
Epoch: [5] [10/60] eta: 0:00:30 lr: 0.000500 loss: 0.1657 (0.1810) loss_classifier: 0.0221 (0.0235) loss_box_reg: 0.0321 (0.0388) loss_mask: 0.1072 (0.1156) loss_ob
jctness: 0.0003 (0.0005) loss_rpn_box_reg: 0.0021 (0.0026) time: 0.6149 data: 0.0409 max mem: 3748
Epoch: [5] [20/60] eta: 0:00:23 lr: 0.000500 loss: 0.1657 (0.1786) loss_classifier: 0.0221 (0.0248) loss_box_reg: 0.0321 (0.0373) loss_mask: 0.1064 (0.1133) loss_ob
jctness: 0.0003 (0.0006) loss_rpn_box_reg: 0.0021 (0.0025) time: 0.5997 data: 0.0365 max mem: 3748
Epoch: [5] [30/60] eta: 0:00:18 lr: 0.000500 loss: 0.1862 (0.1854) loss_classifier: 0.0262 (0.0254) loss_box_reg: 0.0383 (0.0392) loss_mask: 0.1127 (0.1176) loss_ob
jctness: 0.0003 (0.0006) loss_rpn_box_reg: 0.0027 (0.0027) time: 0.5969 data: 0.0366 max mem: 3748
Epoch: [5] [40/60] eta: 0:00:12 lr: 0.000500 loss: 0.1862 (0.1862) loss_classifier: 0.0262 (0.0256) loss_box_reg: 0.0368 (0.0403) loss_mask: 0.1148 (0.1170) loss_ob
jctness: 0.0004 (0.0006) loss_rpn_box_reg: 0.0028 (0.0027) time: 0.6123 data: 0.0406 max mem: 3748
Epoch: [5] [50/60] eta: 0:00:06 lr: 0.000500 loss: 0.1724 (0.1849) loss_classifier: 0.0255 (0.0256) loss_box_reg: 0.0304 (0.0402) loss_mask: 0.1068 (0.1159) loss_ob
jctness: 0.0004 (0.0006) loss_rpn_box_reg: 0.0022 (0.0027) time: 0.6083 data: 0.0412 max mem: 3748
Epoch: [5] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1579 (0.1812) loss_classifier: 0.0213 (0.0250) loss_box_reg: 0.0283 (0.0389) loss_mask: 0.1046 (0.1140) loss_ob
jctness: 0.0004 (0.0006) loss_rpn_box_reg: 0.0021 (0.0027) time: 0.6049 data: 0.0379 max mem: 3748
Epoch: [5] Total time: 0:00:36 (0.6040 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:09 model_time: 0.1698 (0.1698) evaluator_time: 0.0062 (0.0062) time: 0.1938 data: 0.0171 max mem: 3748
Test: [49/50] eta: 0:00:00 model_time: 0.0985 (0.1068) evaluator_time: 0.0040 (0.0067) time: 0.1281 data: 0.0181 max mem: 3748
Test: Total time: 0:00:06 (0.1331 s / it)
Averaged stats: model_time: 0.0985 (0.1068) evaluator_time: 0.0040 (0.0067)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.830
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.988
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.950
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.413
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.652
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.848
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.361
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.859
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.859
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.874
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.752
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.979
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.892
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.385
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.546
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.333
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.782
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.782
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.756
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.790
Model 1 Epoch: 6
Epoch: [6] [ 0/60] eta: 0:00:43 lr: 0.000050 loss: 0.1812 (0.1812) loss_classifier: 0.0300 (0.0300) loss_box_reg: 0.0430 (0.0430) loss_mask: 0.1058 (0.1058) loss_ob
jctness: 0.0002 (0.0002) loss_rpn_box_reg: 0.0022 (0.0022) time: 0.7193 data: 0.0380 max mem: 3748
Epoch: [6] [10/60] eta: 0:00:29 lr: 0.000050 loss: 0.1659 (0.1628) loss_classifier: 0.0232 (0.0218) loss_box_reg: 0.0334 (0.0313) loss_mask: 0.1058 (0.1075) loss_ob
jctness: 0.0002 (0.0005) loss_rpn_box_reg: 0.0014 (0.0016) time: 0.5810 data: 0.0329 max mem: 3748
Epoch: [6] [20/60] eta: 0:00:24 lr: 0.000050 loss: 0.1659 (0.1718) loss_classifier: 0.0220 (0.0237) loss_box_reg: 0.0270 (0.0344) loss_mask: 0.1067 (0.1108) loss_ob
jctness: 0.0002 (0.0007) loss_rpn_box_reg: 0.0021 (0.0023) time: 0.5953 data: 0.0363 max mem: 3748
Epoch: [6] [30/60] eta: 0:00:18 lr: 0.000050 loss: 0.1665 (0.1756) loss_classifier: 0.0245 (0.0252) loss_box_reg: 0.0295 (0.0366) loss_mask: 0.1018 (0.1108) loss_ob
jctness: 0.0003 (0.0006) loss_rpn_box_reg: 0.0021 (0.0024) time: 0.6139 data: 0.0381 max mem: 3748
Epoch: [6] [40/60] eta: 0:00:12 lr: 0.000050 loss: 0.1640 (0.1768) loss_classifier: 0.0218 (0.0244) loss_box_reg: 0.0311 (0.0360) loss_mask: 0.1035 (0.1134) loss_ob
jctness: 0.0003 (0.0005) loss_rpn_box_reg: 0.0020 (0.0025) time: 0.6043 data: 0.0376 max mem: 3748
Epoch: [6] [50/60] eta: 0:00:06 lr: 0.000050 loss: 0.1698 (0.1787) loss_classifier: 0.0218 (0.0248) loss_box_reg: 0.0318 (0.0373) loss_mask: 0.1164 (0.1136) loss_ob
jctness: 0.0003 (0.0006) loss_rpn_box_reg: 0.0019 (0.0024) time: 0.6026 data: 0.0370 max mem: 3748
Epoch: [6] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1795 (0.1817) loss_classifier: 0.0246 (0.0254) loss_box_reg: 0.0394 (0.0381) loss_mask: 0.1164 (0.1148) loss_ob
jctness: 0.0005 (0.0008) loss_rpn_box_reg: 0.0025 (0.0026) time: 0.6173 data: 0.0388 max mem: 3748
Epoch: [6] Total time: 0:00:36 (0.6076 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:09 model_time: 0.1704 (0.1704) evaluator_time: 0.0053 (0.0053) time: 0.1924 data: 0.0160 max mem: 3748
Test: [49/50] eta: 0:00:00 model_time: 0.0999 (0.1063) evaluator_time: 0.0045 (0.0061) time: 0.1340 data: 0.0207 max mem: 3748
Test: Total time: 0:00:06 (0.1314 s / it)
Averaged stats: model_time: 0.0999 (0.1063) evaluator_time: 0.0045 (0.0061)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.829
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.988
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.950
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.413
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.652

```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.847
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.361
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.859
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.859
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.874
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.752
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.978
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.892
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.385
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.546
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.766
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.331
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.781
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.781
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.756
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.789
Model 1 Epoch: 7
Epoch: [7] [ 0/60] eta: 0:00:44 lr: 0.000050 loss: 0.1408 (0.1408) loss_classifier: 0.0232 (0.0232) loss_box_reg: 0.0253 (0.0253) loss_mask: 0.0902 (0.0902) loss_ob
jctness: 0.0004 (0.0004) loss_rpn_box_reg: 0.0017 (0.0017) time: 0.7354 data: 0.0668 max mem: 3748
Epoch: [7] [10/60] eta: 0:00:30 lr: 0.000050 loss: 0.1746 (0.1848) loss_classifier: 0.0262 (0.0266) loss_box_reg: 0.0324 (0.0415) loss_mask: 0.1022 (0.1133) loss_ob
jctness: 0.0003 (0.0007) loss_rpn_box_reg: 0.0017 (0.0026) time: 0.6174 data: 0.0494 max mem: 3748
Epoch: [7] [20/60] eta: 0:00:24 lr: 0.000050 loss: 0.1710 (0.1811) loss_classifier: 0.0247 (0.0256) loss_box_reg: 0.0276 (0.0371) loss_mask: 0.1077 (0.1152) loss_ob
jctness: 0.0003 (0.0006) loss_rpn_box_reg: 0.0017 (0.0026) time: 0.6016 data: 0.0415 max mem: 3748
Epoch: [7] [30/60] eta: 0:00:18 lr: 0.000050 loss: 0.1710 (0.1803) loss_classifier: 0.0246 (0.0247) loss_box_reg: 0.0276 (0.0382) loss_mask: 0.1095 (0.1138) loss_ob
jctness: 0.0004 (0.0010) loss_rpn_box_reg: 0.0023 (0.0027) time: 0.5976 data: 0.0362 max mem: 3748
Epoch: [7] [40/60] eta: 0:00:12 lr: 0.000050 loss: 0.1876 (0.1803) loss_classifier: 0.0246 (0.0251) loss_box_reg: 0.0425 (0.0383) loss_mask: 0.1071 (0.1133) loss_ob
jctness: 0.0005 (0.0010) loss_rpn_box_reg: 0.0024 (0.0026) time: 0.6119 data: 0.0368 max mem: 3748
Epoch: [7] [50/60] eta: 0:00:06 lr: 0.000050 loss: 0.1731 (0.1800) loss_classifier: 0.0230 (0.0252) loss_box_reg: 0.0323 (0.0372) loss_mask: 0.1144 (0.1143) loss_ob
jctness: 0.0003 (0.0009) loss_rpn_box_reg: 0.0020 (0.0025) time: 0.6093 data: 0.0351 max mem: 3748
Epoch: [7] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1665 (0.1781) loss_classifier: 0.0218 (0.0247) loss_box_reg: 0.0281 (0.0369) loss_mask: 0.1071 (0.1130) loss_ob
jctness: 0.0003 (0.0009) loss_rpn_box_reg: 0.0021 (0.0026) time: 0.5948 data: 0.0325 max mem: 3748
Epoch: [7] Total time: 0:00:36 (0.6048 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:09 model_time: 0.1646 (0.1646) evaluator_time: 0.0053 (0.0053) time: 0.1884 data: 0.0178 max mem: 3748
Test: [49/50] eta: 0:00:00 model_time: 0.0976 (0.1067) evaluator_time: 0.0040 (0.0063) time: 0.1302 data: 0.0191 max mem: 3748
Test: Total time: 0:00:06 (0.1327 s / it)
Averaged stats: model_time: 0.0976 (0.1067) evaluator_time: 0.0040 (0.0063)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.829
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.988
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.950
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.413
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.641
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.847
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.361
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.859
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.859
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.874
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.754
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.978
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.887
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.385
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.546
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.769
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.333
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.784
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.784
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.756
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.792
Model 1 Epoch: 8
Epoch: [8] [ 0/60] eta: 0:00:31 lr: 0.000050 loss: 0.1698 (0.1698) loss_classifier: 0.0236 (0.0236) loss_box_reg: 0.0266 (0.0266) loss_mask: 0.1169 (0.1169) loss_ob
jctness: 0.0003 (0.0003) loss_rpn_box_reg: 0.0024 (0.0024) time: 0.5304 data: 0.0425 max mem: 3748
Epoch: [8] [10/60] eta: 0:00:29 lr: 0.000050 loss: 0.1698 (0.1759) loss_classifier: 0.0217 (0.0220) loss_box_reg: 0.0330 (0.0352) loss_mask: 0.1169 (0.1159) loss_ob
jctness: 0.0003 (0.0008) loss_rpn_box_reg: 0.0018 (0.0020) time: 0.5902 data: 0.0359 max mem: 3748
Epoch: [8] [20/60] eta: 0:00:24 lr: 0.000050 loss: 0.1806 (0.1819) loss_classifier: 0.0219 (0.0245) loss_box_reg: 0.0330 (0.0391) loss_mask: 0.1092 (0.1150) loss_ob
jctness: 0.0003 (0.0010) loss_rpn_box_reg: 0.0018 (0.0024) time: 0.6041 data: 0.0398 max mem: 3748
Epoch: [8] [30/60] eta: 0:00:18 lr: 0.000050 loss: 0.1670 (0.1746) loss_classifier: 0.0226 (0.0235) loss_box_reg: 0.0266 (0.0363) loss_mask: 0.1026 (0.1115) loss_ob
jctness: 0.0002 (0.0009) loss_rpn_box_reg: 0.0025 (0.0024) time: 0.6140 data: 0.0411 max mem: 3748
Epoch: [8] [40/60] eta: 0:00:12 lr: 0.000050 loss: 0.1483 (0.1742) loss_classifier: 0.0201 (0.0235) loss_box_reg: 0.0257 (0.0365) loss_mask: 0.1026 (0.1111) loss_ob
jctness: 0.0002 (0.0008) loss_rpn_box_reg: 0.0017 (0.0024) time: 0.6188 data: 0.0392 max mem: 3748
Epoch: [8] [50/60] eta: 0:00:06 lr: 0.000050 loss: 0.1573 (0.1716) loss_classifier: 0.0160 (0.0228) loss_box_reg: 0.0240 (0.0349) loss_mask: 0.1080 (0.1108) loss_ob
jctness: 0.0002 (0.0007) loss_rpn_box_reg: 0.0018 (0.0024) time: 0.5940 data: 0.0368 max mem: 3748
Epoch: [8] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1598 (0.1730) loss_classifier: 0.0234 (0.0240) loss_box_reg: 0.0267 (0.0351) loss_mask: 0.1055 (0.1107) loss_ob
jctness: 0.0004 (0.0008) loss_rpn_box_reg: 0.0021 (0.0025) time: 0.5995 data: 0.0352 max mem: 4110
Epoch: [8] Total time: 0:00:36 (0.6078 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:09 model_time: 0.1705 (0.1705) evaluator_time: 0.0090 (0.0090) time: 0.1971 data: 0.0168 max mem: 4110
Test: [49/50] eta: 0:00:00 model_time: 0.0985 (0.1068) evaluator_time: 0.0040 (0.0052) time: 0.1266 data: 0.0168 max mem: 4110
Test: Total time: 0:00:06 (0.1301 s / it)
Averaged stats: model_time: 0.0985 (0.1068) evaluator_time: 0.0040 (0.0052)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.829
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.988
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.950
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.413
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.641
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.847
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.361
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.860
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.860
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.875
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.754
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.978
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.899

```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.385
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.539
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.769
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.331
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.783
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.783
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.744
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.791
Model 1 Epoch: 9
Epoch: [9] [ 0/60] eta: 0:00:36 lr: 0.000005 loss: 0.1867 (0.1867) loss_classifier: 0.0263 (0.0263) loss_box_reg: 0.0442 (0.0442) loss_mask: 0.1138 (0.1138) loss_obj
jectness: 0.0003 (0.0003) loss_rpn_box_reg: 0.0021 (0.0021) time: 0.6014 data: 0.0356 max mem: 4110
Epoch: [9] [10/60] eta: 0:00:30 lr: 0.000005 loss: 0.1539 (0.1587) loss_classifier: 0.0241 (0.0209) loss_box_reg: 0.0286 (0.0302) loss_mask: 0.1004 (0.1047) loss_obj
jectness: 0.0004 (0.0008) loss_rpn_box_reg: 0.0020 (0.0021) time: 0.6057 data: 0.0414 max mem: 4110
Epoch: [9] [20/60] eta: 0:00:23 lr: 0.000005 loss: 0.1513 (0.1632) loss_classifier: 0.0174 (0.0212) loss_box_reg: 0.0232 (0.0310) loss_mask: 0.1020 (0.1082) loss_obj
jectness: 0.0003 (0.0006) loss_rpn_box_reg: 0.0019 (0.0021) time: 0.5890 data: 0.0369 max mem: 4110
Epoch: [9] [30/60] eta: 0:00:17 lr: 0.000005 loss: 0.1691 (0.1729) loss_classifier: 0.0248 (0.0228) loss_box_reg: 0.0368 (0.0339) loss_mask: 0.1126 (0.1133) loss_obj
jectness: 0.0004 (0.0007) loss_rpn_box_reg: 0.0023 (0.0023) time: 0.5811 data: 0.0346 max mem: 4110
Epoch: [9] [40/60] eta: 0:00:12 lr: 0.000005 loss: 0.1996 (0.1778) loss_classifier: 0.0282 (0.0243) loss_box_reg: 0.0393 (0.0364) loss_mask: 0.1126 (0.1139) loss_obj
jectness: 0.0004 (0.0007) loss_rpn_box_reg: 0.0023 (0.0025) time: 0.6141 data: 0.0396 max mem: 4110
Epoch: [9] [50/60] eta: 0:00:05 lr: 0.000005 loss: 0.1686 (0.1771) loss_classifier: 0.0267 (0.0235) loss_box_reg: 0.0333 (0.0359) loss_mask: 0.1075 (0.1146) loss_obj
jectness: 0.0004 (0.0008) loss_rpn_box_reg: 0.0020 (0.0023) time: 0.6066 data: 0.0383 max mem: 4110
Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.1673 (0.1787) loss_classifier: 0.0214 (0.0241) loss_box_reg: 0.0308 (0.0371) loss_mask: 0.1094 (0.1142) loss_obj
jectness: 0.0005 (0.0008) loss_rpn_box_reg: 0.0021 (0.0025) time: 0.6094 data: 0.0357 max mem: 4110
Epoch: [9] Total time: 0:00:36 (0.6017 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:09 model_time: 0.1697 (0.1697) evaluator_time: 0.0051 (0.0051) time: 0.1918 data: 0.0162 max mem: 4110
Test: [49/50] eta: 0:00:00 model_time: 0.0994 (0.1072) evaluator_time: 0.0041 (0.0064) time: 0.1346 data: 0.0209 max mem: 4110
Test: Total time: 0:00:06 (0.1336 s / it)
Averaged stats: model_time: 0.0994 (0.1072) evaluator_time: 0.0041 (0.0064)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.829
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.988
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.950
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.413
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.641
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.847
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.361
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.860
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.860
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.767
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.875
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.751
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.978
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.899
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.385
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.539
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.766
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.330
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.780
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.780
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.744
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.789

```

That's it! Model 1

In [14]: `#MODEL_OPTION2`

```

model_option2 = get_model_instance_segmentation_option2(num_classes)
model_option2.to(device)

# construct an optimizer for option2
params = [p for p in model_option2.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(params, lr=0.005,
momentum=0.9, weight_decay=0.0005)

# and a learning rate scheduler for the same
lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
step_size=3,
gamma=0.1)

print ("MODEL OPTION 2")

for epoch in range(num_epochs):

    print("Model 2 Epoch: {}".format(epoch))

    # train for one epoch, printing every 10 iterations
    train_one_epoch(model_option2, optimizer, data_loader, device, epoch, print_freq=10)

    # update the learning rate
    lr_scheduler.step()

    # evaluate on the test dataset
    evaluate(model_option2, data_loader_test, device=device)

print ("\n")
print("That's it! Model 2")

```

Downloading: "https://download.pytorch.org/models/mobilenet_v2-7ebf99e0.pth" to /root/.cache/torch/hub/checkpoints/mobilenet_v2-7ebf99e0.pth
0%| | 0.00/13.6M [00:00<?, ?B/s]

```

MODEL OPTION 2
Model 2 Epoch: 0
Epoch: [0] [ 0/60] eta: 0:00:36 lr: 0.000090 loss: 1.5483 (1.5483) loss_classifier: 0.7352 (0.7352) loss_box_reg: 0.0793 (0.0793) loss_objectness: 0.7081 (0.7081) l
oss_rpn_box_reg: 0.0257 (0.0257) time: 0.6098 data: 0.0252 max mem: 4110
Epoch: [0] [10/60] eta: 0:00:20 lr: 0.000936 loss: 1.4805 (1.4755) loss_classifier: 0.6844 (0.6521) loss_box_reg: 0.0793 (0.0757) loss_objectness: 0.6987 (0.6954) l
oss_rpn_box_reg: 0.0506 (0.0523) time: 0.4007 data: 0.0353 max mem: 5714
Epoch: [0] [20/60] eta: 0:00:15 lr: 0.001783 loss: 1.1994 (1.2825) loss_classifier: 0.4023 (0.4896) loss_box_reg: 0.0807 (0.0973) loss_objectness: 0.6620 (0.6526) l
oss_rpn_box_reg: 0.0349 (0.0430) time: 0.3801 data: 0.0346 max mem: 6171
Epoch: [0] [30/60] eta: 0:00:11 lr: 0.002629 loss: 1.0186 (1.1833) loss_classifier: 0.2883 (0.4310) loss_box_reg: 0.1316 (0.1218) loss_objectness: 0.5236 (0.5883) l
oss_rpn_box_reg: 0.0349 (0.0421) time: 0.3802 data: 0.0377 max mem: 6171
Epoch: [0] [40/60] eta: 0:00:07 lr: 0.003476 loss: 0.8882 (1.1042) loss_classifier: 0.2730 (0.3972) loss_box_reg: 0.1691 (0.1437) loss_objectness: 0.3652 (0.5223) l
oss_rpn_box_reg: 0.0381 (0.0410) time: 0.3797 data: 0.0381 max mem: 6171
Epoch: [0] [50/60] eta: 0:00:03 lr: 0.004323 loss: 0.7166 (1.0127) loss_classifier: 0.2349 (0.3635) loss_box_reg: 0.1691 (0.1471) loss_objectness: 0.2655 (0.4636) l
oss_rpn_box_reg: 0.0324 (0.0384) time: 0.3751 data: 0.0342 max mem: 6171
Epoch: [0] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.5928 (0.9590) loss_classifier: 0.2266 (0.3439) loss_box_reg: 0.1759 (0.1565) loss_objectness: 0.1994 (0.4216) l
oss_rpn_box_reg: 0.0231 (0.0370) time: 0.3756 data: 0.0379 max mem: 6171
Epoch: [0] Total time: 0:00:22 (0.3823 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:04 model_time: 0.0546 (0.0546) evaluator_time: 0.0107 (0.0107) time: 0.0823 data: 0.0163 max mem: 6171
Test: [49/50] eta: 0:00:00 model_time: 0.0328 (0.0359) evaluator_time: 0.0025 (0.0035) time: 0.0537 data: 0.0166 max mem: 6171
Test: Total time: 0:00:02 (0.0566 s / it)
Averaged stats: model_time: 0.0328 (0.0359) evaluator_time: 0.0025 (0.0035)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.033
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.105
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.004
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.097
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.036
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.208
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.337
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.033
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.370
Model 2 Epoch: 1
Epoch: [1] [ 0/60] eta: 0:00:22 lr: 0.005000 loss: 0.5127 (0.5127) loss_classifier: 0.1759 (0.1759) loss_box_reg: 0.1739 (0.1739) loss_objectness: 0.1374 (0.1374) l
oss_rpn_box_reg: 0.0256 (0.0256) time: 0.3745 data: 0.0383 max mem: 6171
Epoch: [1] [10/60] eta: 0:00:18 lr: 0.005000 loss: 0.5127 (0.5067) loss_classifier: 0.1759 (0.1746) loss_box_reg: 0.1603 (0.1513) loss_objectness: 0.1390 (0.1538) l
oss_rpn_box_reg: 0.0284 (0.0270) time: 0.3690 data: 0.0334 max mem: 6171
Epoch: [1] [20/60] eta: 0:00:14 lr: 0.005000 loss: 0.4856 (0.5248) loss_classifier: 0.1535 (0.1770) loss_box_reg: 0.1603 (0.1709) loss_objectness: 0.1360 (0.1477) l
oss_rpn_box_reg: 0.0284 (0.0292) time: 0.3737 data: 0.0333 max mem: 6171
Epoch: [1] [30/60] eta: 0:00:11 lr: 0.005000 loss: 0.4856 (0.5197) loss_classifier: 0.1565 (0.1767) loss_box_reg: 0.1639 (0.1730) loss_objectness: 0.1248 (0.1423) l
oss_rpn_box_reg: 0.0244 (0.0277) time: 0.3826 data: 0.0362 max mem: 6171
Epoch: [1] [40/60] eta: 0:00:07 lr: 0.005000 loss: 0.4909 (0.5159) loss_classifier: 0.1580 (0.1708) loss_box_reg: 0.1813 (0.1795) loss_objectness: 0.1248 (0.1374) l
oss_rpn_box_reg: 0.0235 (0.0282) time: 0.3854 data: 0.0378 max mem: 6171
Epoch: [1] [50/60] eta: 0:00:03 lr: 0.005000 loss: 0.4838 (0.5030) loss_classifier: 0.1422 (0.1662) loss_box_reg: 0.1787 (0.1787) loss_objectness: 0.1128 (0.1309) l
oss_rpn_box_reg: 0.0235 (0.0273) time: 0.3765 data: 0.0347 max mem: 6171
Epoch: [1] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.4595 (0.5034) loss_classifier: 0.1332 (0.1642) loss_box_reg: 0.1787 (0.1828) loss_objectness: 0.1061 (0.1283) l
oss_rpn_box_reg: 0.0236 (0.0281) time: 0.3814 data: 0.0405 max mem: 6171
Epoch: [1] Total time: 0:00:22 (0.3798 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:04 model_time: 0.0600 (0.0600) evaluator_time: 0.0039 (0.0039) time: 0.0827 data: 0.0181 max mem: 6171
Test: [49/50] eta: 0:00:00 model_time: 0.0363 (0.0379) evaluator_time: 0.0020 (0.0027) time: 0.0566 data: 0.0168 max mem: 6171
Test: Total time: 0:00:02 (0.0585 s / it)
Averaged stats: model_time: 0.0363 (0.0379) evaluator_time: 0.0020 (0.0027)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.158
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.510
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.055
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.176
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.116
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.339
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.370
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.409
Model 2 Epoch: 2
Epoch: [2] [ 0/60] eta: 0:00:24 lr: 0.005000 loss: 0.5657 (0.5657) loss_classifier: 0.1616 (0.1616) loss_box_reg: 0.2584 (0.2584) loss_objectness: 0.1075 (0.1075) l
oss_rpn_box_reg: 0.0383 (0.0383) time: 0.4106 data: 0.0463 max mem: 6171
Epoch: [2] [10/60] eta: 0:00:19 lr: 0.005000 loss: 0.3937 (0.4336) loss_classifier: 0.1302 (0.1313) loss_box_reg: 0.1717 (0.1799) loss_objectness: 0.0883 (0.0938) l
oss_rpn_box_reg: 0.0226 (0.0286) time: 0.3803 data: 0.0364 max mem: 6171
Epoch: [2] [20/60] eta: 0:00:15 lr: 0.005000 loss: 0.3937 (0.4440) loss_classifier: 0.1302 (0.1346) loss_box_reg: 0.1717 (0.1912) loss_objectness: 0.0815 (0.0914) l
oss_rpn_box_reg: 0.0189 (0.0266) time: 0.3802 data: 0.0395 max mem: 6171
Epoch: [2] [30/60] eta: 0:00:11 lr: 0.005000 loss: 0.4298 (0.4288) loss_classifier: 0.1327 (0.1300) loss_box_reg: 0.2017 (0.1866) loss_objectness: 0.0814 (0.0868) l
oss_rpn_box_reg: 0.0234 (0.0255) time: 0.3804 data: 0.0395 max mem: 6264
Epoch: [2] [40/60] eta: 0:00:07 lr: 0.005000 loss: 0.4020 (0.4208) loss_classifier: 0.1207 (0.1266) loss_box_reg: 0.1513 (0.1819) loss_objectness: 0.0774 (0.0865) l
oss_rpn_box_reg: 0.0253 (0.0258) time: 0.3727 data: 0.0346 max mem: 6264
Epoch: [2] [50/60] eta: 0:00:03 lr: 0.005000 loss: 0.3682 (0.4048) loss_classifier: 0.1056 (0.1216) loss_box_reg: 0.1464 (0.1741) loss_objectness: 0.0764 (0.0835) l
oss_rpn_box_reg: 0.0227 (0.0255) time: 0.3695 data: 0.0342 max mem: 6264
Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.3682 (0.4099) loss_classifier: 0.1056 (0.1238) loss_box_reg: 0.1556 (0.1793) loss_objectness: 0.0668 (0.0812) l
oss_rpn_box_reg: 0.0215 (0.0255) time: 0.3782 data: 0.0391 max mem: 6264
Epoch: [2] Total time: 0:00:22 (0.3782 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:04 model_time: 0.0564 (0.0564) evaluator_time: 0.0048 (0.0048) time: 0.0909 data: 0.0288 max mem: 6264
Test: [49/50] eta: 0:00:00 model_time: 0.0338 (0.0363) evaluator_time: 0.0037 (0.0043) time: 0.0618 data: 0.0224 max mem: 6264
Test: Total time: 0:00:03 (0.0639 s / it)
Averaged stats: model_time: 0.0338 (0.0363) evaluator_time: 0.0037 (0.0043)
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.206
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.552
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.063
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.228
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.141
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.381
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.412
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= medium | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.455
Model 2 Epoch: 3
Epoch: [3] [ 0/60] eta: 0:00:22 lr: 0.000500 loss: 0.4481 (0.4481) loss_classifier: 0.1347 (0.1347) loss_box_reg: 0.2129 (0.2129) loss_objectness: 0.0810 (0.0810) l
oss_rpn_box_reg: 0.0195 (0.0195) time: 0.3754 data: 0.0440 max mem: 6264

```

```

Epoch: [3] [10/60] eta: 0:00:19 lr: 0.000500 loss: 0.3439 (0.3628) loss_classifier: 0.1153 (0.1113) loss_box_reg: 0.1754 (0.1707) loss_objectness: 0.0569 (0.0573) 1
oss_rpn_box_reg: 0.0234 (0.0235) time: 0.3869 data: 0.0380 max mem: 6264
Epoch: [3] [20/60] eta: 0:00:15 lr: 0.000500 loss: 0.3216 (0.3389) loss_classifier: 0.0992 (0.1041) loss_box_reg: 0.1535 (0.1563) loss_objectness: 0.0542 (0.0577) 1
oss_rpn_box_reg: 0.0201 (0.0208) time: 0.3821 data: 0.0383 max mem: 6264
Epoch: [3] [30/60] eta: 0:00:11 lr: 0.000500 loss: 0.3216 (0.3517) loss_classifier: 0.1030 (0.1079) loss_box_reg: 0.1612 (0.1644) loss_objectness: 0.0571 (0.0585) 1
oss_rpn_box_reg: 0.0184 (0.0209) time: 0.3786 data: 0.0372 max mem: 6264
Epoch: [3] [40/60] eta: 0:00:07 lr: 0.000500 loss: 0.3549 (0.3523) loss_classifier: 0.1067 (0.1080) loss_box_reg: 0.1705 (0.1640) loss_objectness: 0.0620 (0.0580) 1
oss_rpn_box_reg: 0.0228 (0.0224) time: 0.3842 data: 0.0362 max mem: 6264
Epoch: [3] [50/60] eta: 0:00:03 lr: 0.000500 loss: 0.3358 (0.3526) loss_classifier: 0.1067 (0.1090) loss_box_reg: 0.1589 (0.1649) loss_objectness: 0.0577 (0.0572) 1
oss_rpn_box_reg: 0.0208 (0.0215) time: 0.3863 data: 0.0402 max mem: 6264
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.3065 (0.3492) loss_classifier: 0.1011 (0.1070) loss_box_reg: 0.1562 (0.1631) loss_objectness: 0.0480 (0.0581) 1
oss_rpn_box_reg: 0.0183 (0.0211) time: 0.3807 data: 0.0394 max mem: 6264
Epoch: [3] Total time: 0:00:22 (0.3819 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:03 model_time: 0.0543 (0.0543) evaluator_time: 0.0025 (0.0025) time: 0.0744 data: 0.0169 max mem: 6264
Test: [49/50] eta: 0:00:00 model_time: 0.0326 (0.0352) evaluator_time: 0.0017 (0.0019) time: 0.0515 data: 0.0161 max mem: 6264
Test: Total time: 0:00:02 (0.0540 s / it)
Averaged stats: model_time: 0.0326 (0.0352) evaluator_time: 0.0017 (0.0019)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.222
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.599
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.084
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.246
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.157
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.386
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.414
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.457
Model 2 Epoch: 4
Epoch: [4] [ 0/60] eta: 0:00:23 lr: 0.000500 loss: 0.3060 (0.3060) loss_classifier: 0.0976 (0.0976) loss_box_reg: 0.1501 (0.1501) loss_objectness: 0.0435 (0.0435) 1
oss_rpn_box_reg: 0.0148 (0.0148) time: 0.3883 data: 0.0459 max mem: 6264
Epoch: [4] [10/60] eta: 0:00:19 lr: 0.000500 loss: 0.3060 (0.3268) loss_classifier: 0.0976 (0.1039) loss_box_reg: 0.1409 (0.1477) loss_objectness: 0.0590 (0.0565) 1
oss_rpn_box_reg: 0.0148 (0.0188) time: 0.3977 data: 0.0474 max mem: 6264
Epoch: [4] [20/60] eta: 0:00:15 lr: 0.000500 loss: 0.3483 (0.3647) loss_classifier: 0.1080 (0.1158) loss_box_reg: 0.1570 (0.1722) loss_objectness: 0.0538 (0.0567) 1
oss_rpn_box_reg: 0.0173 (0.0201) time: 0.3888 data: 0.0413 max mem: 6264
Epoch: [4] [30/60] eta: 0:00:11 lr: 0.000500 loss: 0.3649 (0.3614) loss_classifier: 0.1056 (0.1116) loss_box_reg: 0.1827 (0.1712) loss_objectness: 0.0523 (0.0576) 1
oss_rpn_box_reg: 0.0194 (0.0211) time: 0.3780 data: 0.0349 max mem: 6264
Epoch: [4] [40/60] eta: 0:00:07 lr: 0.000500 loss: 0.3387 (0.3632) loss_classifier: 0.0948 (0.1108) loss_box_reg: 0.1673 (0.1717) loss_objectness: 0.0498 (0.0586) 1
oss_rpn_box_reg: 0.0232 (0.0222) time: 0.3841 data: 0.0390 max mem: 6264
Epoch: [4] [50/60] eta: 0:00:03 lr: 0.000500 loss: 0.3387 (0.3586) loss_classifier: 0.1016 (0.1096) loss_box_reg: 0.1673 (0.1700) loss_objectness: 0.0504 (0.0577) 1
oss_rpn_box_reg: 0.0192 (0.0213) time: 0.3847 data: 0.0403 max mem: 6264
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.2903 (0.3500) loss_classifier: 0.0981 (0.1075) loss_box_reg: 0.1418 (0.1655) loss_objectness: 0.0504 (0.0563) 1
oss_rpn_box_reg: 0.0150 (0.0207) time: 0.3776 data: 0.0356 max mem: 6264
Epoch: [4] Total time: 0:00:22 (0.3829 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:03 model_time: 0.0543 (0.0543) evaluator_time: 0.0024 (0.0024) time: 0.0760 data: 0.0186 max mem: 6264
Test: [49/50] eta: 0:00:00 model_time: 0.0343 (0.0363) evaluator_time: 0.0033 (0.0050) time: 0.0662 data: 0.0219 max mem: 6264
Test: Total time: 0:00:03 (0.0621 s / it)
Averaged stats: model_time: 0.0343 (0.0363) evaluator_time: 0.0033 (0.0050)
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.279
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.645
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.157
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.002
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.308
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.182
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.429
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.454
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.067
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.497
Model 2 Epoch: 5
Epoch: [5] [ 0/60] eta: 0:00:23 lr: 0.000500 loss: 0.6110 (0.6110) loss_classifier: 0.1749 (0.1749) loss_box_reg: 0.2841 (0.2841) loss_objectness: 0.1050 (0.1050) 1
oss_rpn_box_reg: 0.0470 (0.0470) time: 0.3990 data: 0.0475 max mem: 6264
Epoch: [5] [10/60] eta: 0:00:18 lr: 0.000500 loss: 0.3202 (0.3145) loss_classifier: 0.1059 (0.0987) loss_box_reg: 0.1407 (0.1428) loss_objectness: 0.0445 (0.0554) 1
oss_rpn_box_reg: 0.0142 (0.0177) time: 0.3774 data: 0.0368 max mem: 6264
Epoch: [5] [20/60] eta: 0:00:15 lr: 0.000500 loss: 0.3202 (0.3502) loss_classifier: 0.1026 (0.1043) loss_box_reg: 0.1407 (0.1639) loss_objectness: 0.0498 (0.0616) 1
oss_rpn_box_reg: 0.0161 (0.0204) time: 0.3790 data: 0.0359 max mem: 6264
Epoch: [5] [30/60] eta: 0:00:11 lr: 0.000500 loss: 0.3115 (0.3456) loss_classifier: 0.1009 (0.1048) loss_box_reg: 0.1383 (0.1634) loss_objectness: 0.0521 (0.0579) 1
oss_rpn_box_reg: 0.0183 (0.0196) time: 0.3765 data: 0.0356 max mem: 6264
Epoch: [5] [40/60] eta: 0:00:07 lr: 0.000500 loss: 0.3016 (0.3376) loss_classifier: 0.0947 (0.1021) loss_box_reg: 0.1300 (0.1581) loss_objectness: 0.0525 (0.0580) 1
oss_rpn_box_reg: 0.0164 (0.0194) time: 0.3746 data: 0.0389 max mem: 6264
Epoch: [5] [50/60] eta: 0:00:03 lr: 0.000500 loss: 0.3181 (0.3369) loss_classifier: 0.0858 (0.1008) loss_box_reg: 0.1355 (0.1588) loss_objectness: 0.0574 (0.0569) 1
oss_rpn_box_reg: 0.0168 (0.0203) time: 0.3805 data: 0.0384 max mem: 6264
Epoch: [5] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.3293 (0.3430) loss_classifier: 0.0954 (0.1028) loss_box_reg: 0.1463 (0.1619) loss_objectness: 0.0472 (0.0579) 1
oss_rpn_box_reg: 0.0162 (0.0204) time: 0.3846 data: 0.0346 max mem: 6698
Epoch: [5] Total time: 0:00:22 (0.3801 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:04 model_time: 0.0558 (0.0558) evaluator_time: 0.0051 (0.0051) time: 0.0865 data: 0.0248 max mem: 6698
Test: [49/50] eta: 0:00:00 model_time: 0.0344 (0.0365) evaluator_time: 0.0020 (0.0034) time: 0.0565 data: 0.0183 max mem: 6698
Test: Total time: 0:00:03 (0.0614 s / it)
Averaged stats: model_time: 0.0344 (0.0365) evaluator_time: 0.0020 (0.0034)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.244
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.621
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.106
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.004
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.269
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.159
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.394
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.418
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.067
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.457
Model 2 Epoch: 6
Epoch: [6] [ 0/60] eta: 0:00:20 lr: 0.000050 loss: 0.3441 (0.3441) loss_classifier: 0.1117 (0.1117) loss_box_reg: 0.1667 (0.1667) loss_objectness: 0.0453 (0.0453) 1
oss_rpn_box_reg: 0.0204 (0.0204) time: 0.3436 data: 0.0268 max mem: 6698
Epoch: [6] [10/60] eta: 0:00:18 lr: 0.000050 loss: 0.3446 (0.3596) loss_classifier: 0.1105 (0.1085) loss_box_reg: 0.1865 (0.1809) loss_objectness: 0.0481 (0.0506) 1
oss_rpn_box_reg: 0.0204 (0.0197) time: 0.3729 data: 0.0367 max mem: 6698
Epoch: [6] [20/60] eta: 0:00:15 lr: 0.000050 loss: 0.2979 (0.3221) loss_classifier: 0.0930 (0.0967) loss_box_reg: 0.1443 (0.1573) loss_objectness: 0.0454 (0.0483) 1
oss_rpn_box_reg: 0.0179 (0.0197) time: 0.3775 data: 0.0367 max mem: 6698

```

```

Epoch: [6] [30/60] eta: 0:00:11 lr: 0.000050 loss: 0.2815 (0.3187) loss_classifier: 0.0869 (0.0985) loss_box_reg: 0.1379 (0.1533) loss_objectness: 0.0386 (0.0475) 1
oss_rpn_box_reg: 0.0167 (0.0194) time: 0.3818 data: 0.0380 max mem: 6698
Epoch: [6] [40/60] eta: 0:00:07 lr: 0.000050 loss: 0.2880 (0.3140) loss_classifier: 0.0921 (0.0972) loss_box_reg: 0.1379 (0.1516) loss_objectness: 0.0386 (0.0470) 1
oss_rpn_box_reg: 0.0147 (0.0183) time: 0.3778 data: 0.0371 max mem: 6698
Epoch: [6] [50/60] eta: 0:00:03 lr: 0.000050 loss: 0.2981 (0.3280) loss_classifier: 0.0908 (0.1001) loss_box_reg: 0.1379 (0.1578) loss_objectness: 0.0450 (0.0508) 1
oss_rpn_box_reg: 0.0191 (0.0193) time: 0.3813 data: 0.0365 max mem: 6698
Epoch: [6] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.3217 (0.3301) loss_classifier: 0.1031 (0.1014) loss_box_reg: 0.1474 (0.1591) loss_objectness: 0.0502 (0.0502) 1
oss_rpn_box_reg: 0.0192 (0.0194) time: 0.3892 data: 0.0359 max mem: 6698
Epoch: [6] Total time: 0:00:22 (0.3816 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:03 model_time: 0.0550 (0.0550) evaluator_time: 0.0026 (0.0026) time: 0.0791 data: 0.0207 max mem: 6698
Test: [49/50] eta: 0:00:00 model_time: 0.0347 (0.0363) evaluator_time: 0.0018 (0.0021) time: 0.0548 data: 0.0173 max mem: 6698
Test: Total time: 0:00:02 (0.0565 s / it)
Averaged stats: model_time: 0.0347 (0.0363) evaluator_time: 0.0018 (0.0021)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.248
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.654
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.074
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.003
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.273
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.156
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.404
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.432
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.056
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.473
Model 2 Epoch: 7
Epoch: [7] [ 0/60] eta: 0:00:21 lr: 0.000050 loss: 0.2458 (0.2458) loss_classifier: 0.0645 (0.0645) loss_box_reg: 0.1382 (0.1382) loss_objectness: 0.0258 (0.0258) 1
oss_rpn_box_reg: 0.0173 (0.0173) time: 0.3662 data: 0.0361 max mem: 6698
Epoch: [7] [10/60] eta: 0:00:18 lr: 0.000050 loss: 0.2629 (0.3183) loss_classifier: 0.0960 (0.0969) loss_box_reg: 0.1356 (0.1473) loss_objectness: 0.0418 (0.0530) 1
oss_rpn_box_reg: 0.0176 (0.0211) time: 0.3693 data: 0.0323 max mem: 6698
Epoch: [7] [20/60] eta: 0:00:15 lr: 0.000050 loss: 0.3018 (0.3290) loss_classifier: 0.0960 (0.1004) loss_box_reg: 0.1356 (0.1556) loss_objectness: 0.0425 (0.0525) 1
oss_rpn_box_reg: 0.0190 (0.0206) time: 0.3700 data: 0.0356 max mem: 6698
Epoch: [7] [30/60] eta: 0:00:11 lr: 0.000050 loss: 0.3247 (0.3496) loss_classifier: 0.0976 (0.1062) loss_box_reg: 0.1588 (0.1684) loss_objectness: 0.0499 (0.0544) 1
oss_rpn_box_reg: 0.0190 (0.0207) time: 0.3895 data: 0.0439 max mem: 6698
Epoch: [7] [40/60] eta: 0:00:07 lr: 0.000050 loss: 0.3315 (0.3473) loss_classifier: 0.0988 (0.1055) loss_box_reg: 0.1758 (0.1672) loss_objectness: 0.0500 (0.0539) 1
oss_rpn_box_reg: 0.0204 (0.0207) time: 0.3878 data: 0.0421 max mem: 6698
Epoch: [7] [50/60] eta: 0:00:03 lr: 0.000050 loss: 0.3005 (0.3331) loss_classifier: 0.0882 (0.1012) loss_box_reg: 0.1401 (0.1601) loss_objectness: 0.0481 (0.0524) 1
oss_rpn_box_reg: 0.0173 (0.0194) time: 0.3751 data: 0.0333 max mem: 6698
Epoch: [7] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.2915 (0.3343) loss_classifier: 0.0870 (0.1021) loss_box_reg: 0.1539 (0.1606) loss_objectness: 0.0485 (0.0523) 1
oss_rpn_box_reg: 0.0156 (0.0194) time: 0.3839 data: 0.0375 max mem: 6698
Epoch: [7] Total time: 0:00:22 (0.3824 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:03 model_time: 0.0543 (0.0543) evaluator_time: 0.0025 (0.0025) time: 0.0749 data: 0.0173 max mem: 6698
Test: [49/50] eta: 0:00:00 model_time: 0.0332 (0.0362) evaluator_time: 0.0018 (0.0022) time: 0.0539 data: 0.0171 max mem: 6698
Test: Total time: 0:00:02 (0.0563 s / it)
Averaged stats: model_time: 0.0332 (0.0362) evaluator_time: 0.0018 (0.0022)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.268
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.640
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.166
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.004
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.296
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.181
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.433
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.459
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.056
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.503
Model 2 Epoch: 8
Epoch: [8] [ 0/60] eta: 0:00:22 lr: 0.000050 loss: 0.4425 (0.4425) loss_classifier: 0.1317 (0.1317) loss_box_reg: 0.1979 (0.1979) loss_objectness: 0.0853 (0.0853) 1
oss_rpn_box_reg: 0.0276 (0.0276) time: 0.3714 data: 0.0305 max mem: 6698
Epoch: [8] [10/60] eta: 0:00:18 lr: 0.000050 loss: 0.3537 (0.3539) loss_classifier: 0.1093 (0.1074) loss_box_reg: 0.1728 (0.1715) loss_objectness: 0.0553 (0.0527) 1
oss_rpn_box_reg: 0.0218 (0.0222) time: 0.3745 data: 0.0329 max mem: 6698
Epoch: [8] [20/60] eta: 0:00:15 lr: 0.000050 loss: 0.3457 (0.3768) loss_classifier: 0.1065 (0.1140) loss_box_reg: 0.1633 (0.1859) loss_objectness: 0.0543 (0.0545) 1
oss_rpn_box_reg: 0.0191 (0.0225) time: 0.3882 data: 0.0393 max mem: 6698
Epoch: [8] [30/60] eta: 0:00:11 lr: 0.000050 loss: 0.2902 (0.3482) loss_classifier: 0.0924 (0.1072) loss_box_reg: 0.1319 (0.1685) loss_objectness: 0.0415 (0.0518) 1
oss_rpn_box_reg: 0.0168 (0.0207) time: 0.3912 data: 0.0410 max mem: 6698
Epoch: [8] [40/60] eta: 0:00:07 lr: 0.000050 loss: 0.2577 (0.3306) loss_classifier: 0.0835 (0.1037) loss_box_reg: 0.1242 (0.1621) loss_objectness: 0.0408 (0.0526) 1
oss_rpn_box_reg: 0.0162 (0.0202) time: 0.3834 data: 0.0374 max mem: 6698
Epoch: [8] [50/60] eta: 0:00:03 lr: 0.000050 loss: 0.2839 (0.3280) loss_classifier: 0.0816 (0.1001) loss_box_reg: 0.1317 (0.1564) loss_objectness: 0.0482 (0.0519) 1
oss_rpn_box_reg: 0.0156 (0.0195) time: 0.3818 data: 0.0397 max mem: 6698
Epoch: [8] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.2839 (0.3250) loss_classifier: 0.0880 (0.0991) loss_box_reg: 0.1423 (0.1543) loss_objectness: 0.0434 (0.0526) 1
oss_rpn_box_reg: 0.0146 (0.0190) time: 0.3831 data: 0.0416 max mem: 6698
Epoch: [8] Total time: 0:00:23 (0.3840 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:03 model_time: 0.0550 (0.0550) evaluator_time: 0.0024 (0.0024) time: 0.0759 data: 0.0178 max mem: 6698
Test: [49/50] eta: 0:00:00 model_time: 0.0337 (0.0360) evaluator_time: 0.0017 (0.0021) time: 0.0529 data: 0.0164 max mem: 6698
Test: Total time: 0:00:02 (0.0555 s / it)
Averaged stats: model_time: 0.0337 (0.0360) evaluator_time: 0.0017 (0.0021)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.274
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.674
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.159
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.002
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.303
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.177
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.402
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.433
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.067
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.473
Model 2 Epoch: 9
Epoch: [9] [ 0/60] eta: 0:00:21 lr: 0.000005 loss: 0.2460 (0.2460) loss_classifier: 0.0772 (0.0772) loss_box_reg: 0.0851 (0.0851) loss_objectness: 0.0695 (0.0695) 1
oss_rpn_box_reg: 0.0143 (0.0143) time: 0.3595 data: 0.0282 max mem: 6698
Epoch: [9] [10/60] eta: 0:00:19 lr: 0.000005 loss: 0.3128 (0.3490) loss_classifier: 0.1043 (0.1090) loss_box_reg: 0.1557 (0.1661) loss_objectness: 0.0507 (0.0528) 1
oss_rpn_box_reg: 0.0184 (0.0211) time: 0.3812 data: 0.0373 max mem: 6698
Epoch: [9] [20/60] eta: 0:00:15 lr: 0.000005 loss: 0.3055 (0.3252) loss_classifier: 0.0996 (0.1000) loss_box_reg: 0.1423 (0.1527) loss_objectness: 0.0504 (0.0517) 1
oss_rpn_box_reg: 0.0183 (0.0208) time: 0.3827 data: 0.0378 max mem: 6698
Epoch: [9] [30/60] eta: 0:00:11 lr: 0.000005 loss: 0.3133 (0.3458) loss_classifier: 0.1029 (0.1053) loss_box_reg: 0.1559 (0.1669) loss_objectness: 0.0531 (0.0533) 1
oss_rpn_box_reg: 0.0151 (0.0202) time: 0.3792 data: 0.0380 max mem: 6698
Epoch: [9] [40/60] eta: 0:00:07 lr: 0.000005 loss: 0.3147 (0.3327) loss_classifier: 0.0928 (0.1006) loss_box_reg: 0.1422 (0.1570) loss_objectness: 0.0547 (0.0553) 1
oss_rpn_box_reg: 0.0164 (0.0197) time: 0.3742 data: 0.0354 max mem: 6698

```

```

Epoch: [9] [50/60] eta: 0:00:03 lr: 0.000005 loss: 0.3147 (0.3339) loss_classifier: 0.0905 (0.1009) loss_box_reg: 0.1416 (0.1595) loss_objectness: 0.0440 (0.0539) 1
oss_rpn_box_reg: 0.0171 (0.0196) time: 0.3785 data: 0.0387 max mem: 6698
Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.3137 (0.3334) loss_classifier: 0.0933 (0.1010) loss_box_reg: 0.1546 (0.1593) loss_objectness: 0.0440 (0.0534) 1
oss_rpn_box_reg: 0.0176 (0.0197) time: 0.3840 data: 0.0408 max mem: 6698
Epoch: [9] Total time: 0:00:22 (0.3805 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:03 model_time: 0.0545 (0.0545) evaluator_time: 0.0024 (0.0024) time: 0.0752 data: 0.0176 max mem: 6698
Test: [49/50] eta: 0:00:00 model_time: 0.0351 (0.0366) evaluator_time: 0.0019 (0.0022) time: 0.0553 data: 0.0174 max mem: 6698
Test: Total time: 0:00:02 (0.0572 s / it)
Averaged stats: model_time: 0.0351 (0.0366) evaluator_time: 0.0019 (0.0022)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.273
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.652
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.135
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.003
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.302
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.184
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.431
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.465
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.089
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.507

```

That's it! Model 2

```

In [15]: import torch
import torchvision
from PIL import Image
import matplotlib.pyplot as plt

url = "https://upload.wikimedia.org/wikipedia/en/4/42/Beatles_-_Abbey_Road.jpg"
img = Image.open(requests.get(url, stream=True).raw)

transform = transforms.Compose([transforms.Resize(256),
                                transforms.CenterCrop(224),
                                transforms.ToTensor(),
                                ])

img = transform(img).to(device)

# Plot the original image
plt.imshow(transforms.ToPILImage()(img))
print ("Original Image")
plt.show()

transform = torchvision.transforms.Compose([
    torchvision.transforms.Resize(256),
    torchvision.transforms.CenterCrop(224),

    torchvision.transforms.Normalize(
        mean=[0.485, 0.456, 0.406],
        std=[0.229, 0.224, 0.225]
    )
])

img = transform(img).to(device)

print("OPTION 1")

# Make a prediction
with torch.no_grad():
    output = model_option1(img.unsqueeze(0))
    print('Prediction for model_option1: \n', output)

# Filter the output to keep only the top scoring boxes
threshold = 0.5
_boxes = output[0]['boxes']

# Draw the boxes on the image
fig, ax = plt.subplots()
ax.imshow(transforms.ToPILImage()(img))

for box in _boxes:
    x1, y1, x2, y2 = box.tolist()
    w, h = x2 - x1, y2 - y1
    rect = plt.Rectangle((x1, y1), w, h, linewidth=1, edgecolor='r', facecolor='none')
    ax.add_patch(rect)
plt.show()

print("OPTION 2")

# Make a prediction
with torch.no_grad():
    output = model_option2(img.unsqueeze(0))
    print('Prediction for model_option2: \n', output)


# Filter the output to keep only the top scoring boxes
threshold = 0.5
_boxes = output[0]['boxes']

# Draw the boxes on the image
fig, ax = plt.subplots()
ax.imshow(transforms.ToPILImage()(img))


for box in _boxes:
    x1, y1, x2, y2 = box.tolist()
    w, h = x2 - x1, y2 - y1
    rect = plt.Rectangle((x1, y1), w, h, linewidth=1, edgecolor='r', facecolor='none')
    ax.add_patch(rect)
plt.show()

Original Image


```



OPTION 1
Prediction for model_option1:
[{'boxes': tensor([], device='cuda:0', size=(0, 4)), 'labels': tensor([], device='cuda:0', dtype=torch.int64), 'scores': tensor([], device='cuda:0'), 'masks': tensor([], device='cuda:0', size=(0, 1, 224, 224))}]



OPTION 2
Prediction for model_option2:
[{'boxes': tensor([[61.7875, 91.8146, 108.8170, 217.9420],
[7.7791, 82.5294, 56.8638, 224.0000],
[176.3702, 120.4357, 223.9798, 224.0000],
[64.7320, 133.9791, 87.6553, 202.6269],
[60.2962, 63.2871, 143.0251, 224.0000],
[17.7155, 75.2567, 99.8058, 223.7091],
[95.4509, 26.1168, 119.9425, 81.9522],
[105.6704, 25.2083, 129.5414, 83.2234],
[9.1040, 102.0064, 196.3956, 219.8465],
[54.6946, 3.4210, 136.3288, 107.8350]]], device='cuda:0'), 'labels': tensor([1, 1, 1, 1, 1, 1, 1, 1, 1, 1], device='cuda:0'), 'scores': tensor([0.3437, 0.2554, 0.1789, 0.1712, 0.1502, 0.1249, 0.1001, 0.0676, 0.0614, 0.0581], device='cuda:0')}]



In [15]: